



Introduction to Machine Learning

ENEE 6583 Neural Nets

Dr. Alsamman

Slide Credits: A Ng,



Related Fields

❖ Statistical Estimation

➤ old; math; small; foundation

❖ Pattern Recognition

➤ 60s; images; engineering

❖ Machine Learning

➤ 80s; CS

❖ Artificial Intelligence

➤ ML; CS

❖ Data Mining

➤ 90s; data; info theory



Learning

- ❖ Supervised:
 - Labeled training
- ❖ Unsupervised
 - Unlabeled
- ❖ Semisupervised
 - Label deficient
- ❖ Reinforcement
 - Reward actions



Supervised Learning

❖ Goal: Given a labeled training data set, use a learning function to generate a good predictor of the output for any input.

❖ Notation:

- \mathbf{X} ALL input
- \mathbb{R}^n n = dimension of input: $\mathbf{x} = \{x_1, x_2, \dots, x_n\}$
- \mathbf{Y} ALL output
- $(x^{(i)}, y^{(i)})$ Training sample i
- m Size of training set: $\{(x^{(i)}, y^{(i)}); i = 1, \dots, m\}$
- $h(\mathbf{x}^{(i)})$ Learning function, aka, hypothesis
- $h: \mathbf{X} \rightarrow \mathbf{Y}$ Learning is the mapping of input to output



Linear Regression

- ❖ Linear: line fitting
- ❖ Regression: reduction
- ❖ Learning function:

$$\hat{Y} = h(\mathbf{X}^{(i)}) = \mathbf{W}^T \mathbf{X} + b = w_1 x_1^{(i)} + w_2 x_2^{(i)} + \dots + w_n x_n^{(i)} + b$$

- \mathbf{W}, b : learning parameters
 - \mathbf{X} : input data (aka features)
 - \hat{Y} : prediction
- ❖ Goal: $h(\mathbf{X})$ must be as close to \mathbf{Y} (ground truth, aka labels) as possible




Example

Size in feet ² (x)	Price (\$) in 1000's (y)
2104	460
1416	232
1534	315
852	178
...	...

m=47

- ❖ How many inputs/outputs?
- ❖ Dimension of input/output?
- ❖ Hypothesis function?






Cost Function


- ❖ Measure of how close $h(x^i)$ is to y





- ❖ Least squares cost function:



$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

- ❖ Goal: minimize cost











Gradient Descent

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

- ❖ Gradient: use derivate/slope
- ❖ Descent: make sure that with each iteration cost is decreasing
- ❖ Gradient descent algorithm:
 1. Start with initial values for W, b
 2. Compute J
 3. Update all varaiaables **simultaneously**:

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$


4. Goto 2 until w, b converge.



Gradient Descent

❖ Understanding the effect of gradient descent

- W, b : learning parameters
- Learning rate α : hyperparameter





Gradient Descent

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

$$w = w - \alpha \frac{\partial}{\partial w_j} J(w, b)$$

❖ Batch gradient

$$w = w - \alpha \frac{1}{m} \sum_{i=1}^m ((h(x^{(i)}) - y^{(i)})x^{(i)}), \quad b = b - \alpha \frac{1}{m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})$$

❖ Stochastic Gradient Descent

$$w = w - \alpha (h(x^{(i)}) - y^{(i)})x^{(i)}, \quad b = b - \alpha (h(x^{(i)}) - y^{(i)})$$



Multiple Features: X

Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

❖ $m = 47$

❖ $X = \begin{bmatrix} \mathbf{x}_j^{(i)} \end{bmatrix}$



Multivariate Linear Regression

❖ Multivariate:

$$h(\mathbf{X}^{(i)}) = W_1 x_1^{(i)} + W_2 x_2^{(i)} + \dots + W_n x_n^{(i)} + b$$

$$J(\mathbf{W}) = \frac{1}{2m} \sum_{i=1}^m (h(\mathbf{X}^{(i)}) - \mathbf{Y}^{(i)})^2$$

❖ Gradient Descent:

$$W_j = W_j - \alpha \frac{\partial}{\partial W_j} J(\mathbf{W}, b)$$

$$W_j = W_j - \alpha \frac{1}{m} \sum_{i=1}^m ((h(\mathbf{X}^{(i)}) - \mathbf{Y}^{(i)}) x_j^{(i)})$$

❖ Stochastic Descent:

$$W_j = W_j - \alpha (h(\mathbf{X}^{(i)}) - \mathbf{Y}^{(i)}) x_j^{(i)}$$



Matrix Based Solution

	Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
x_0	x_1	x_2	x_3	x_4	y
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178


$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}, \quad Y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}, \quad W = \begin{bmatrix} W_0 \\ W_1 \\ W_2 \\ W_3 \\ W_4 \end{bmatrix}$$

$$XW = y \Rightarrow w = (X^T X)^{-1} X^T y$$




Univariate Polynomial Regression


$$h(x^{(i)}) = W_1x + W_2x^2 + \dots + W_nx^n + b$$



linear



Order 2



Oder 5



❖ Binary classification

❖ Linear Regression

❖ Logistic Hypothesis

Logistic Regression

$$0 \leq y \leq 1$$

$$\mathbf{x} = [1 \ x_1 \ x_2]$$
$$\mathbf{w} = [b \ w_1 \ w_2]$$

$$h(x) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

$$h(x) = P(y = 1 | \mathbf{x}, \mathbf{w})$$


$$1 - h(x) = P(y = 0 | \mathbf{x}, \mathbf{w})$$

$$P(y | \mathbf{x}; \mathbf{w}) = (h(x))^y (1 - h(x))^{1-y}$$




Sigmoid

$$h(z) = \frac{1}{1 + \exp(-z)}$$





Example: Classify Salmon and Seabass





Likelihood

$$J(\mathbf{W}, b) = \frac{1}{2m} \sum_{i=1}^m (h(\mathbf{X}^{(i)}) - y^{(i)})^2$$

$$L(\mathbf{W}, b) = \prod_{i=1}^m p(y^{(i)} | \mathbf{X}^{(i)}, \mathbf{W}, b) = \prod_{i=1}^m \left(h(\mathbf{X}^{(i)}) \right)^{y^{(i)}} \left(1 - h(\mathbf{X}^{(i)}) \right)^{1-y^{(i)}}$$

❖ Log likelihood:

$$\begin{aligned} \mathcal{L}(\mathbf{W}, b) &= \log(L(\mathbf{W}, b)) \\ &= \sum_{i=1}^m y^{(i)} \log(h(\mathbf{X}^{(i)})) + (1 - y^{(i)}) \log(1 - h(\mathbf{X}^{(i)})) \end{aligned}$$



Gradient Decent

$$w_j = w_j - \frac{\alpha}{m} \frac{\partial}{\partial w_j} l(w)$$

$$\frac{\partial l}{\partial w} = \frac{\partial l}{\partial h} \frac{\partial h}{\partial z} \frac{\partial z}{\partial w}$$



Partial Gradient

$$\frac{\partial l}{\partial h} = \frac{\partial}{\partial h} \left(\sum y \log h + (1 - y) \log(1 - h) \right) = \sum y \frac{\partial}{\partial h} \log h + (1 - y) \frac{\partial}{\partial h} \log(1 - h)$$

$$\frac{d}{dh} \log_a h = \frac{1}{h \ln a} = \frac{1}{h} \Big|_{a=e}$$

$$\frac{d}{dh} \ln 1 - h = \frac{d}{du} \ln u \frac{du}{dh} = -\frac{1}{1 - h}$$

$$\frac{\partial l}{\partial h} = \sum \frac{y}{h} - \frac{(1 - y)}{1 - h}$$



Partial Gradient

$$h(z) = \text{sigmoid}(z) = \frac{1}{1 + e^{-z}}$$

$$h(z) = \frac{1}{u}$$

$$\frac{\partial h}{\partial z} = \frac{\partial h}{\partial u} \frac{\partial u}{\partial z} = \frac{1}{u^2} \frac{du}{dz} = \frac{-e^{-z}}{(1 + e^{-z})^2} = \frac{1}{1 + e^{-z}} \frac{-e^{-z}}{1 + e^{-z}} = \frac{1}{1 + e^{-z}} \left(1 - \frac{1}{1 + e^{-z}} \right)$$

$$\frac{\partial h}{\partial z} = h(1 - h)$$



Gradient

$$z = w^T x$$

$$\frac{\partial z}{\partial w} = x$$

$$\frac{\partial l}{\partial w} = \sum \left(\frac{y}{h} + \frac{(1-y)}{1-h} \right) h(1-h)x = \sum y(1-h)x - (1-y)hx = (y-h)x$$



Gradient Descent

$$w_j = w_j - \frac{\alpha}{m} \frac{\partial}{\partial w_j} l(\mathbf{w})$$



$$w_j = w_j - \frac{\alpha}{m} \sum_{i=1}^m \left(y^{(i)} - h(w_j x_j^{(i)}) \right) x_j^{(i)}$$

❖ Vectorized:

$$\mathbf{W} = \mathbf{W} - \frac{\alpha}{m} \left(\mathbf{X}^T \left(\mathbf{Y} - h(\mathbf{W}^T \mathbf{X}) \right) \right)$$



Multiclass







- ❖ Can treat as a k-binary
 - Good for classes that are non mutually exclusive
- ❖ 1-hot encoding





K-binary

- ❖ Treat 1 class as class '1' and all others as class '0'
- ❖ Repeat for all classes
- ❖ Must train k times
- ❖ Works well when classes are independent




Class 1: 
 Class 2: 
 Class 3: 





1-hot encoded Approach for Sigmoid

- ❖ Trained only once
- ❖ More effective
 - Balanced training set
- ❖ Reformat the GT into a vector
 - K-binary has only 1 output: $Y = 0$ or 1
 - 1-hot encoded output = a vector of k binary values
- ❖ E.g. $k = 3$ classes
 - Y is a vector of 3 values
 - Class 1: $Y = [1 \ 0 \ 0]$
 - Class 2: $Y = [0 \ 1 \ 0]$
 - Class 3: $Y = [0 \ 0 \ 1]$
 - 3 sigmoids sharing inputs





Softmax Regression

- ❖ 1-hot encoded approach
- ❖ New activation function

$$y^{(i)} \in \{1, 2, 3, \dots, M\}$$

$$h(x^{(i)}) = p(y^{(i)} = n | x^{(i)}; \theta) = \frac{\exp(w_n^T x^{(i)})}{\sum_{j=1}^M \exp(w_j^T x^{(i)})} = \frac{\text{class}_n \text{ output}}{\text{Sum of All class outputs}}$$

- ❖ For the class output to be 1
 - Must increase output of class, while suppressing other outputs



Likelihood

$$J = \prod_{i=1}^m \prod_{j=1}^M 1\{y_j^i = n\} h_j$$

$$1\{y_j^i = n\} = \begin{cases} 1 & y_j^i = n \\ 0 & y_j^i \neq n \end{cases}$$

$$L = \log J = -\frac{1}{m} \sum_{i=1}^m \sum_{j=1}^M 1\{y_j^i = n\} \log(h_j) = \frac{1}{m} \sum_{i=1}^m L_j$$

$$L_j = -1\{y_j^i = n\} \log h_j$$



Gradient

$$\frac{\partial L_j}{\partial w_j} = -1\{y_j^i = n\} \frac{\partial}{\partial w_j} \log h_i = -1\{y_j^i = n\} \frac{1}{h_j} \frac{\partial h_j}{\partial w_j}$$

$$\frac{\partial h_j}{\partial w_j} = \begin{cases} h_n(1 - h_j) & j = n \\ -h_j h_n & j \neq n \end{cases}$$

$$\frac{\partial L}{\partial w_j} = (h_j - \{y_j^i = n\})x$$



Stable Softmax

- ❖ Exponential can be numerically unstable
- ❖ use a stabilizing coefficient

$$h(z_i) = \frac{ce^{z_i}}{\sum_j ce^{z_j}} = \frac{e^{z_i+D}}{\sum_j e^{z_j+D}}$$

➤ $D = -\max\{z_j \forall j\}$



Gradient Descent