



Recurrent NN

ENEE 4583/5583 Neural Nets

Dr. Alsamman

Slide Credits:



Sequential Data Input

- ❖ Ordered data
- ❖ Spatial dependent order: text
- ❖ Physical order: chemical, DNA sequence
- ❖ Time order: audio, finance, medical
- ❖ Spatial and time order: video,

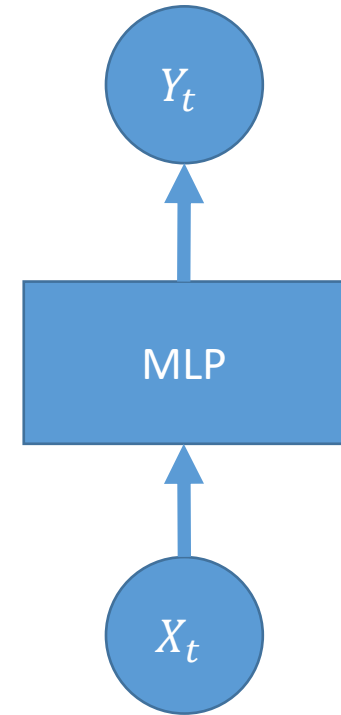
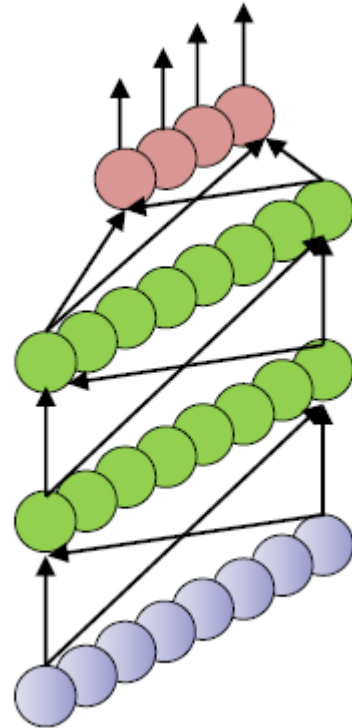


Sequential Driven Output

- ❖ Single output dependent on past sequence of input
 - Classification of a sequence
 - E.g. sentiment classification
- ❖ Sequence output dependent on past single input
 - Sequence generation
 - E.g. image caption
- ❖ sequence output dependent on current sequence of input
 - Updated prediction
 - E.g. stock prediction
- ❖ Current sequence output dependent on past sequence of input
 - Delayed prediction
 - E.g. language translation

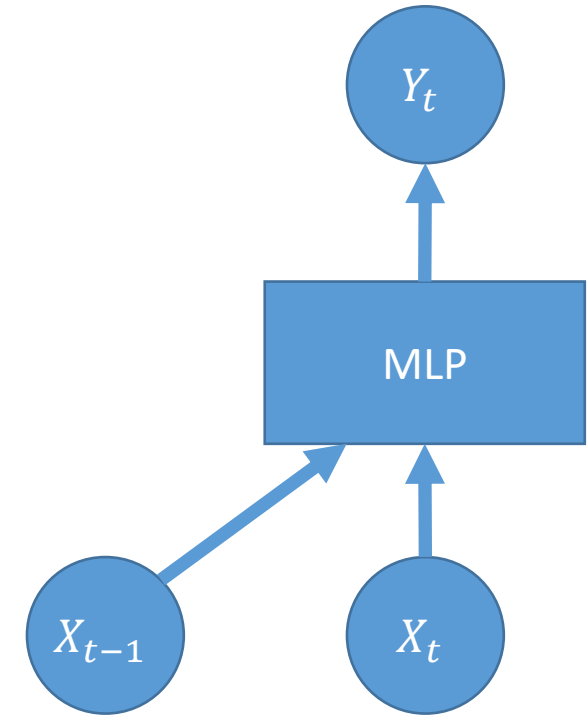
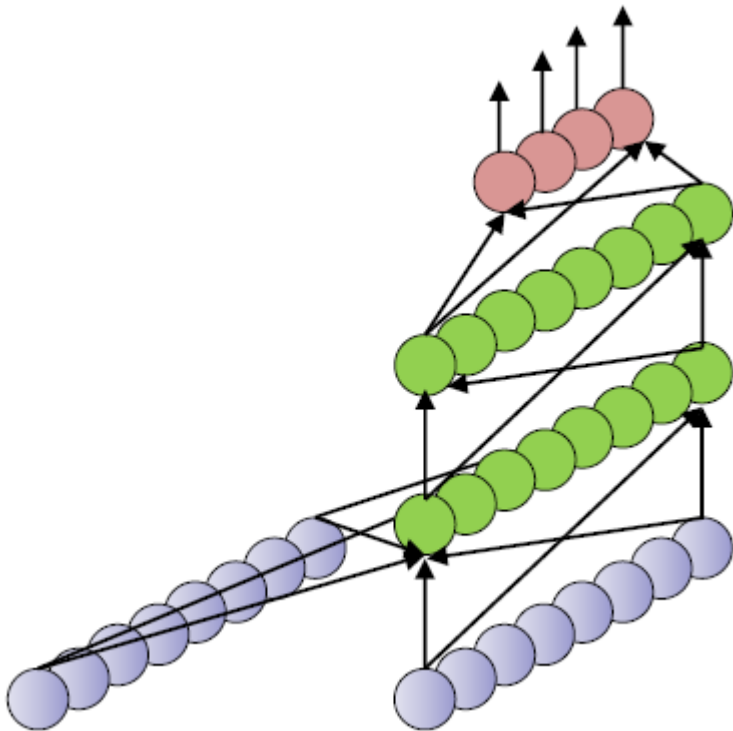


Diagram Representation: Example 1



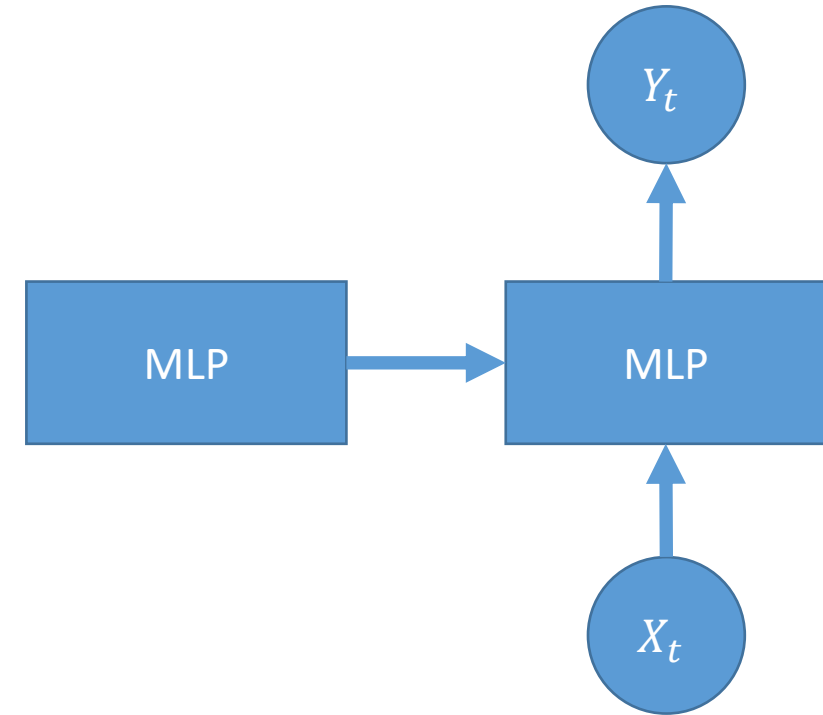
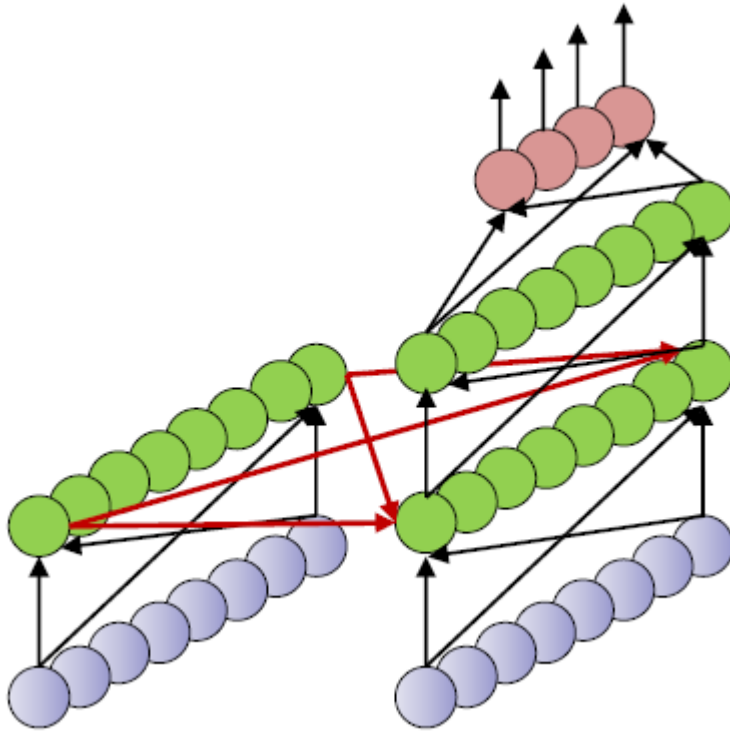


Example 2





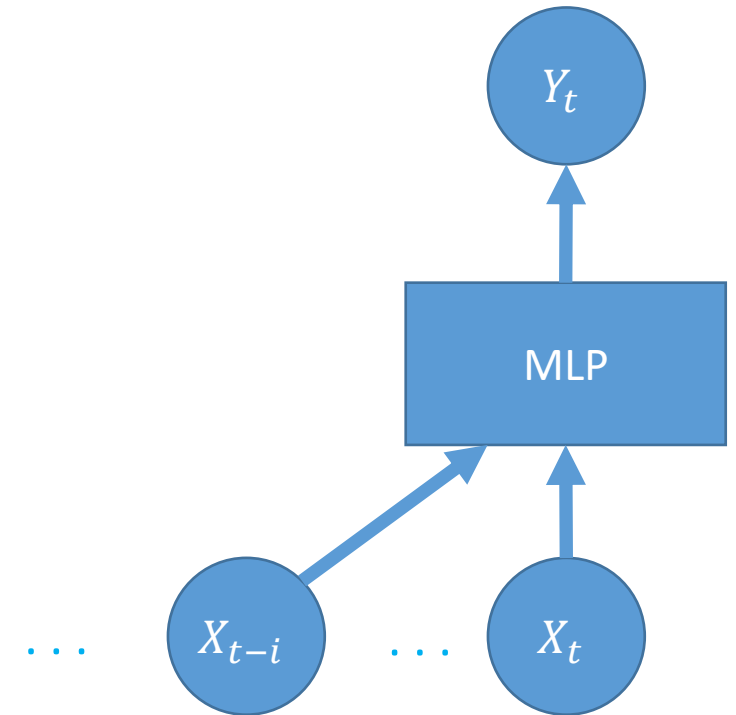
Example 3





Problems with MLP

- ❖ Problem 1: Sequence inputs can be arbitrary length
- ❖ Solution: fixed window size as input
- ❖ Problem 2: choosing problem size
- ❖ Problem 3: number of parameters can explode
 - 100 hidden neurons x 100 inputs x 100 words > 1M





TDNN (1989)

- ❖ Time delayed NN

- E.g stock predictor

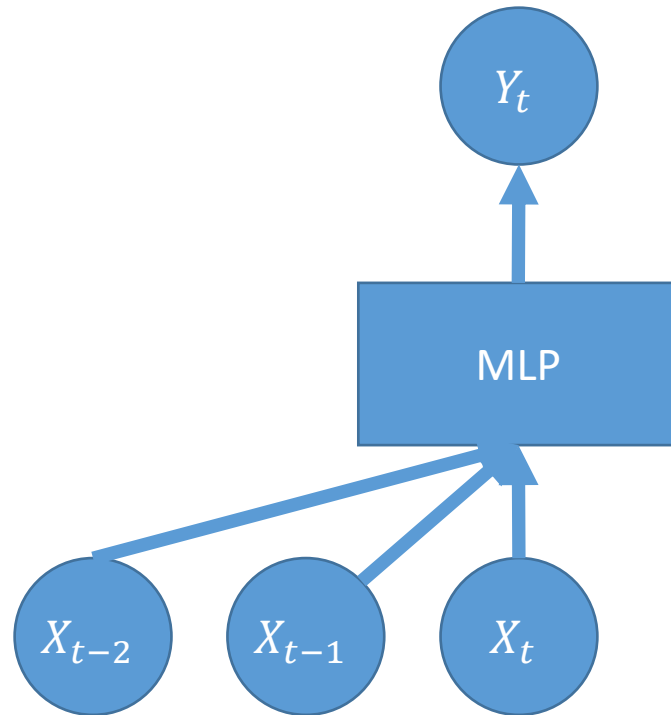
- ❖ Paper: *Phoneme Recognition Using Time-Delay Neural Networks*

- Weibel, Hanzawa, Hinton, Shikano

- Inspired by Fukushima's Neocognitron

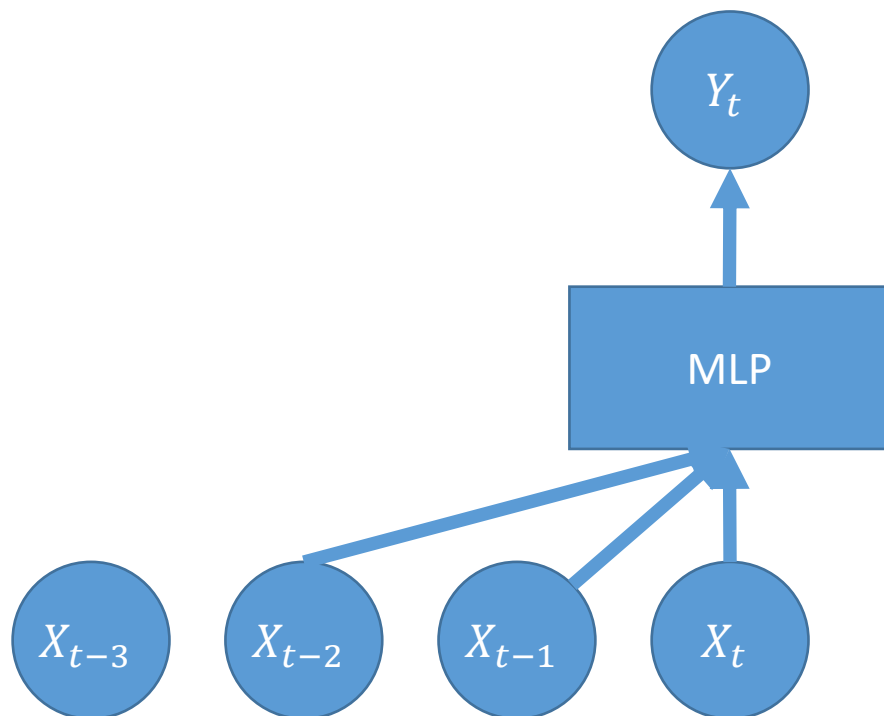


$T=3$



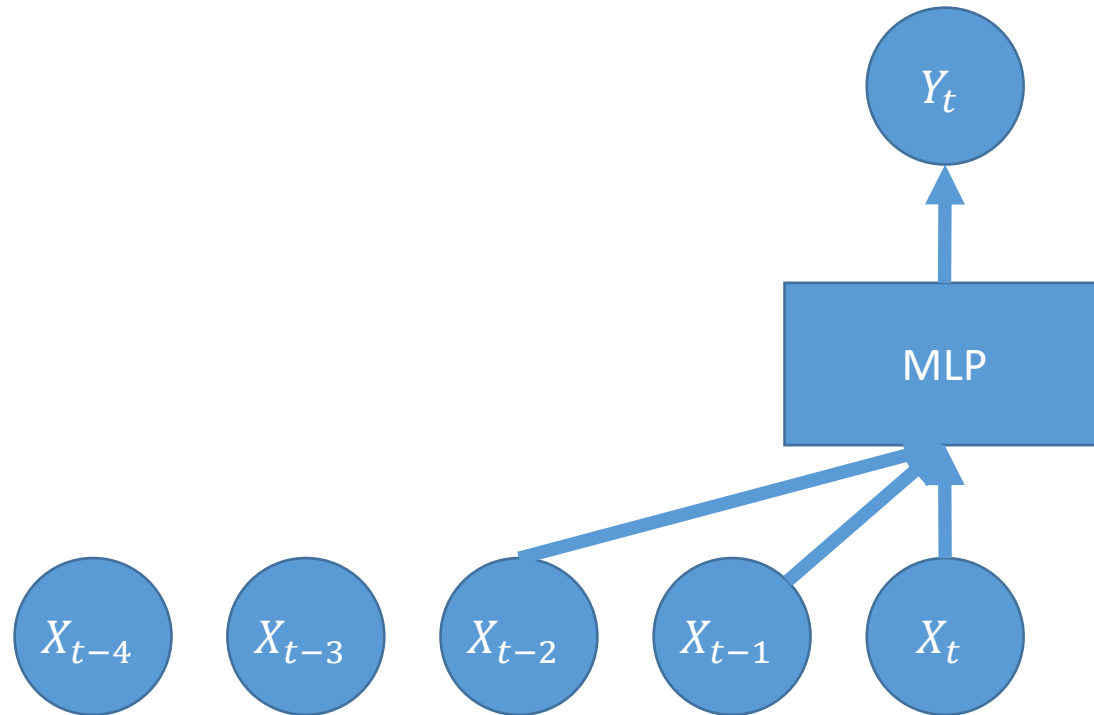


$T=4$



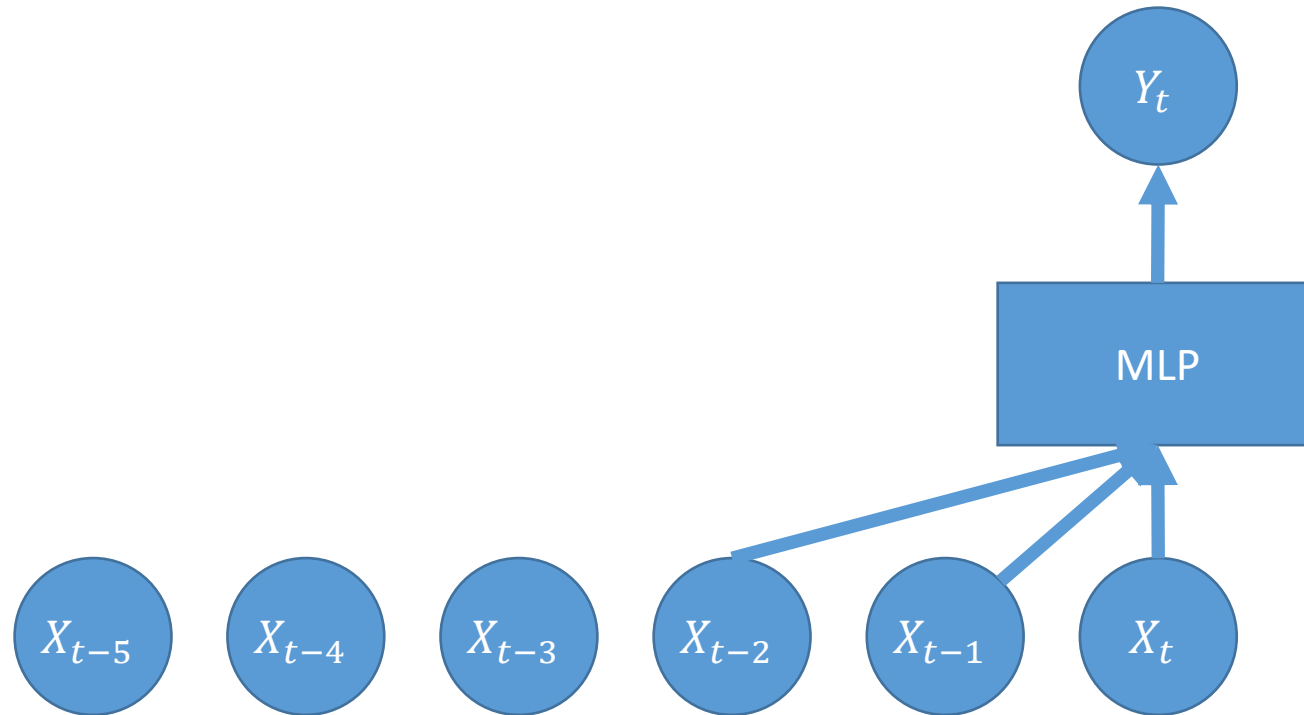


$T=5$





$T=6$





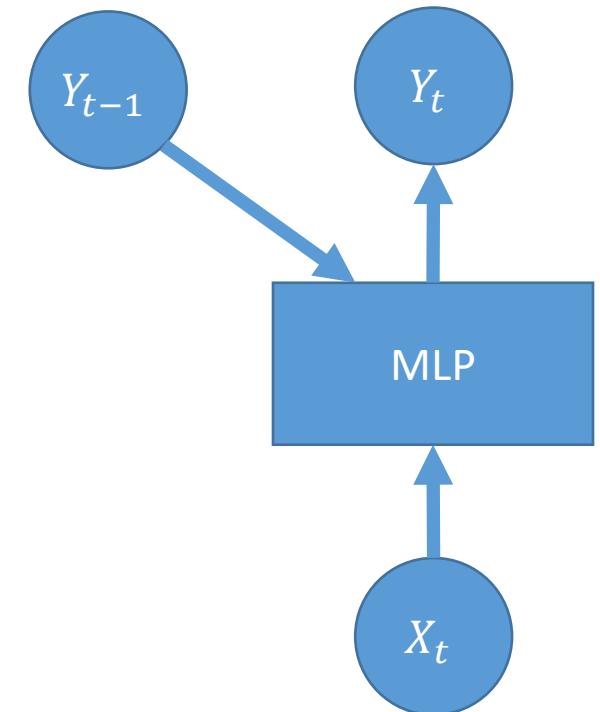
Limitations of TDNN

- ❖ Finite response system:
 - Output driven by past $N-T=s$ only
- ❖ Sliding predictor
 - Much like a convolutional nnet
- ❖ Problem: some trends are “seasonal”
 - Bias the output
- ❖ Prefer: Infinite response system
 - Would like to learn “trends”
 - With weaker and weaker influence



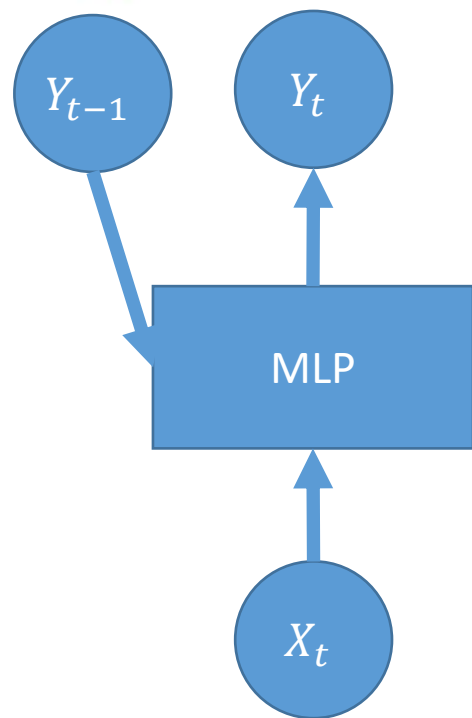
NARX (1985)

- ❖ Nonlinear autoregressive exogenous model
 - Leontaritis & Billings
- ❖ Recursion from output
- ❖ Popular for time-series
 - Weather
 - Stock-market
 - Tracking



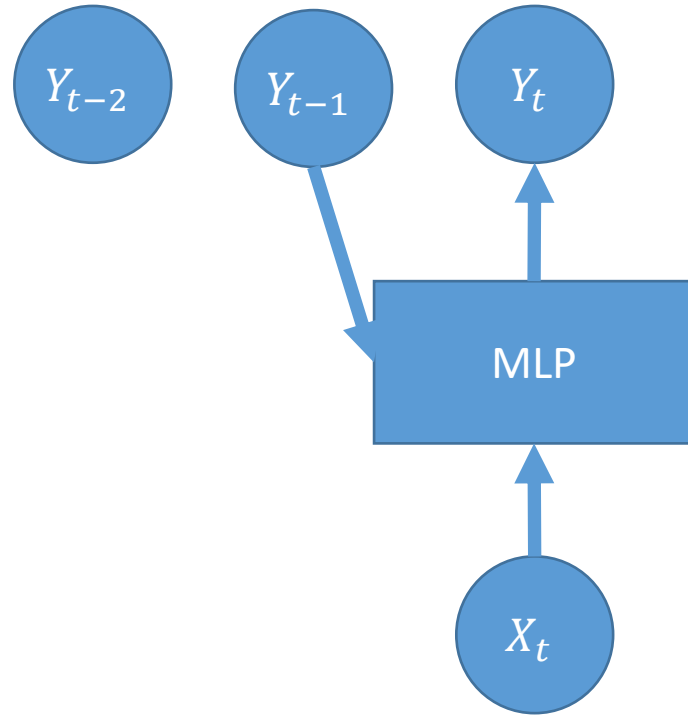


$T=2$



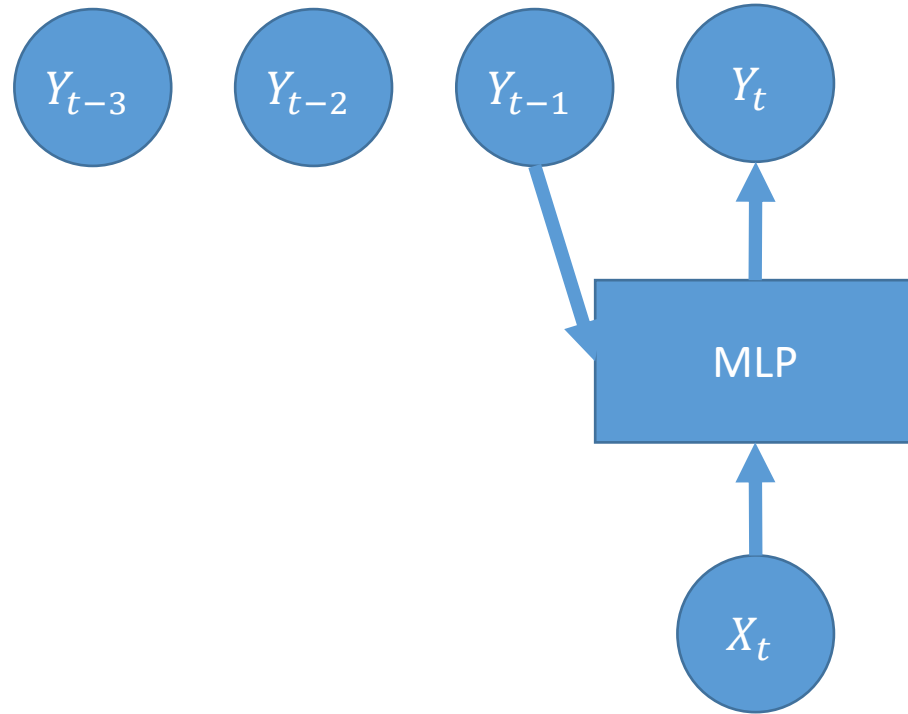


$T=3$



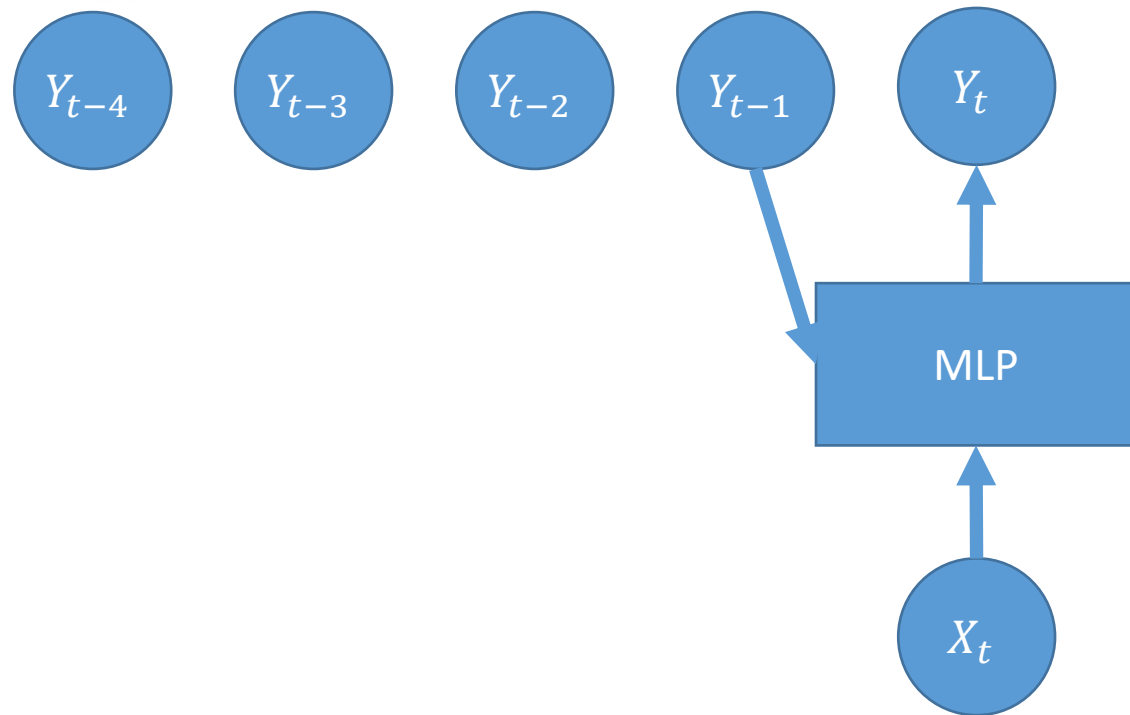


$T=4$





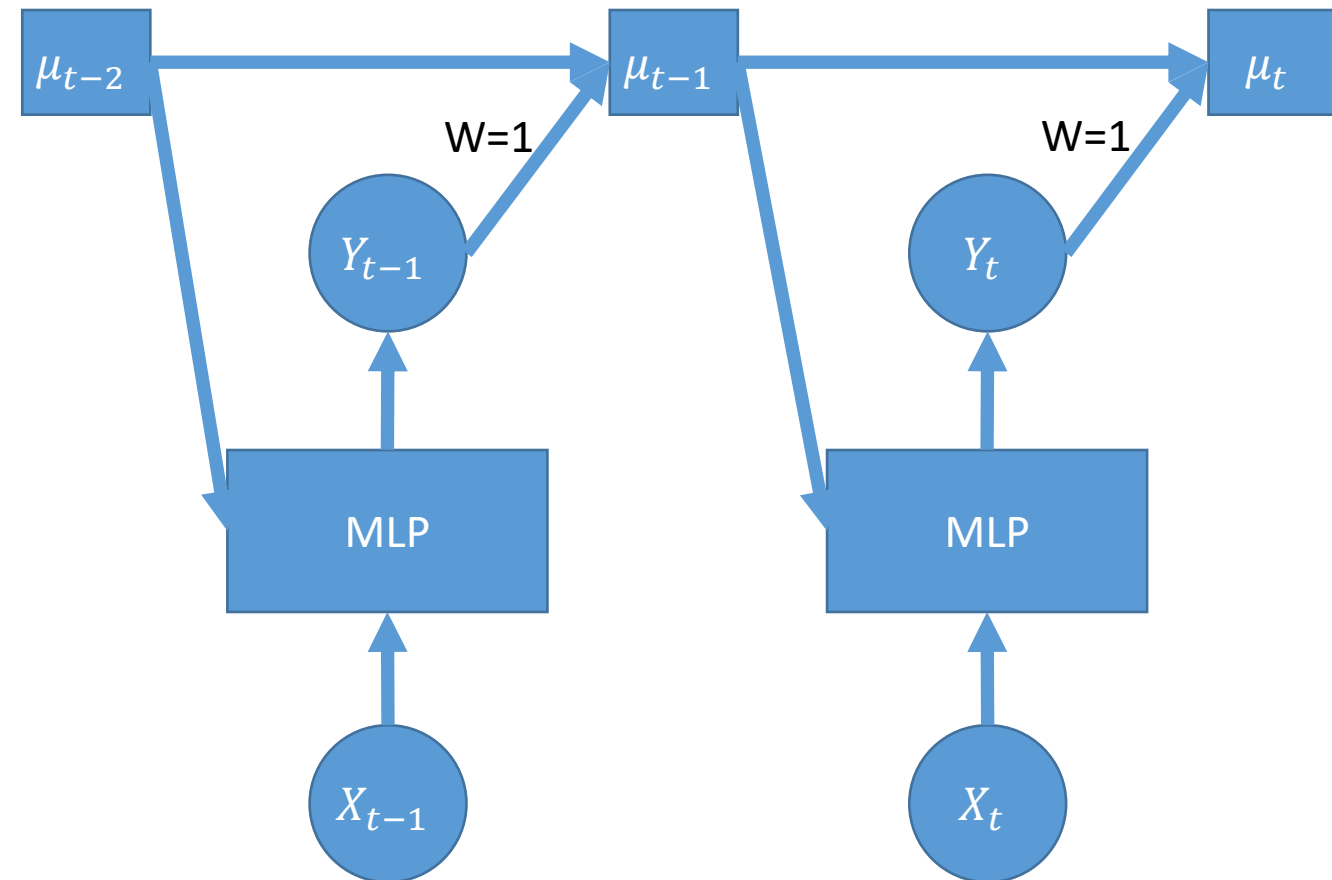
$T=5$





Jordan Net (1986)

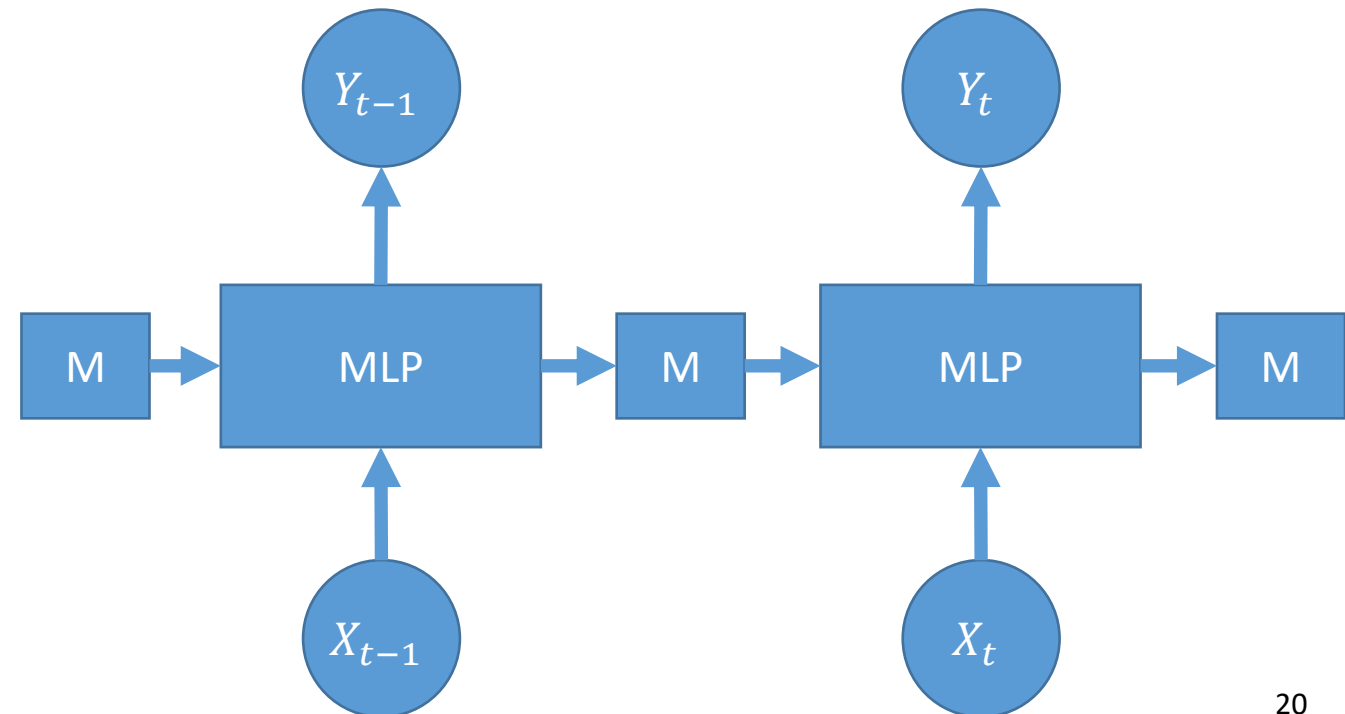
- ❖ Memory is a running average of outputs
 - Stored statistic
 - Doesn't learn to remember
- ❖ Fixed weights
 - $W_{Y\mu} = 1$





Elman Net (1990)

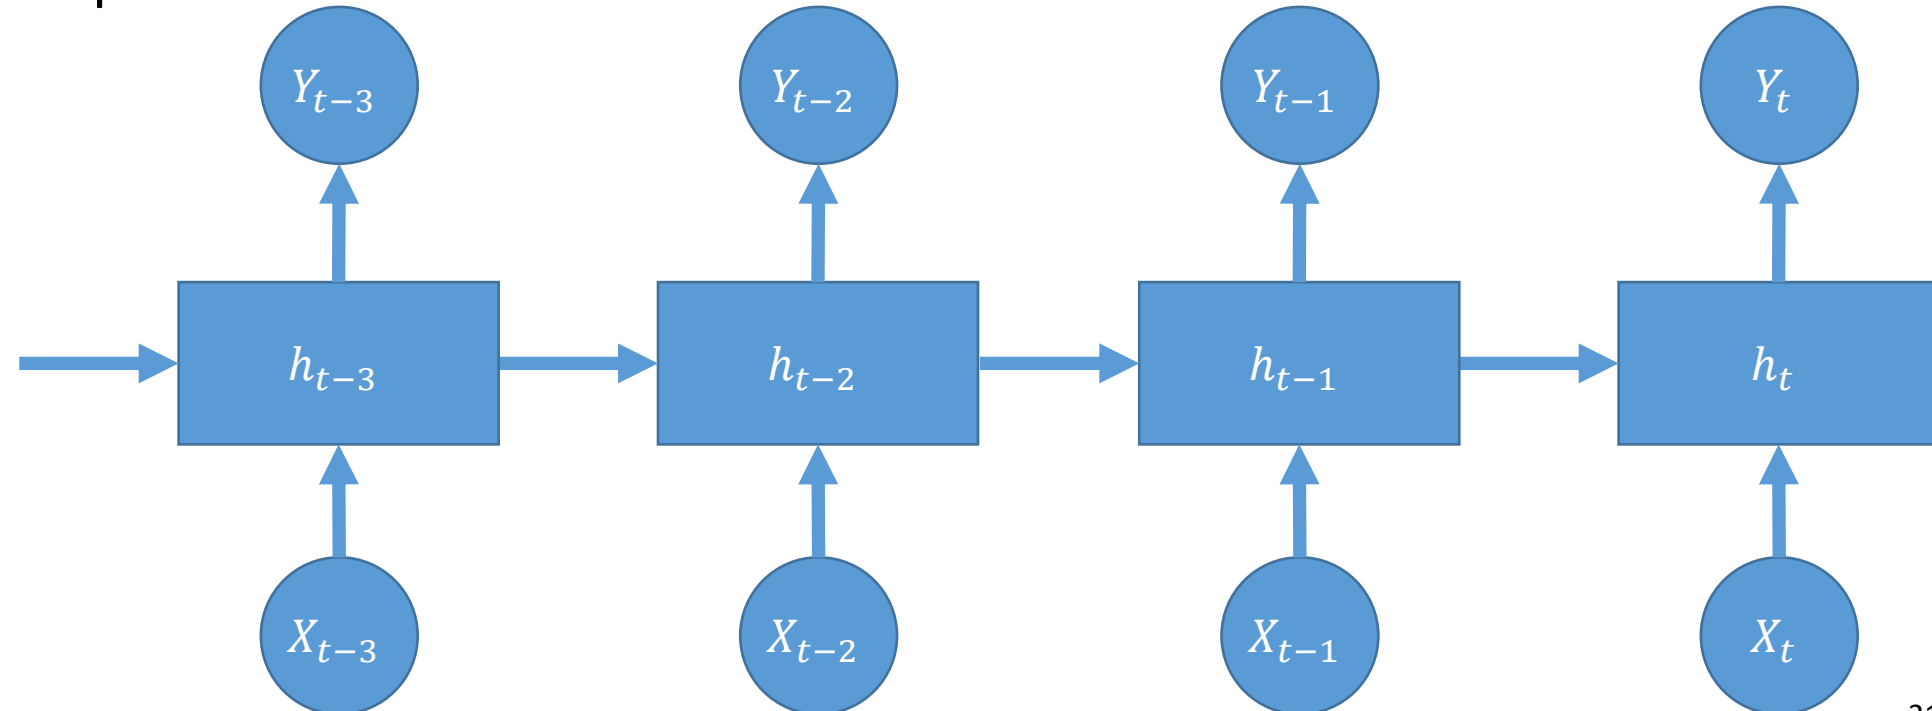
- ❖ M is memory: store the previous state
- ❖ Only the weight from M to MLP is learned
- ❖ M is approximated as independent 1-step history nets
 - Backprop only back 1 step
 - Can't backprop to the beginning





State-Space Model

- ❖ Fix number of input
- ❖ Share parameters of MLP at each step
 - Memory is embedded into state, h
- ❖ MLP arbitrarily complex



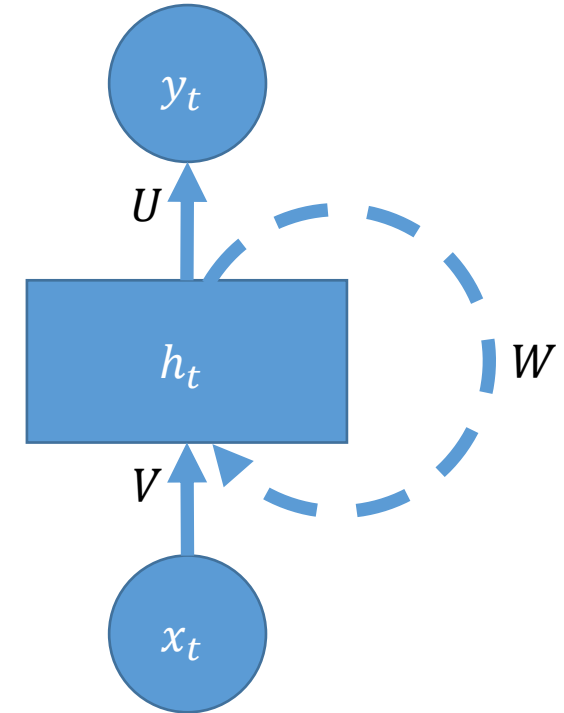


Single Recurrent Network Model

❖ Folded Model (in time)

❖ Parameters

- x_t : input sequence at t
- y_t : output (prediction)
- h : state of network
- V : weights of input
- U : weights of outputs
- W : shared weights
- b_h : biases for hidden state
- b_y : bias for output
- f : typically tanh function
- g : for classification softmax is typically used



$$h_t = f(Vx_t + Wh_{t-1} + b_h)$$

$$y_t = g(Uh_t + b_y)$$



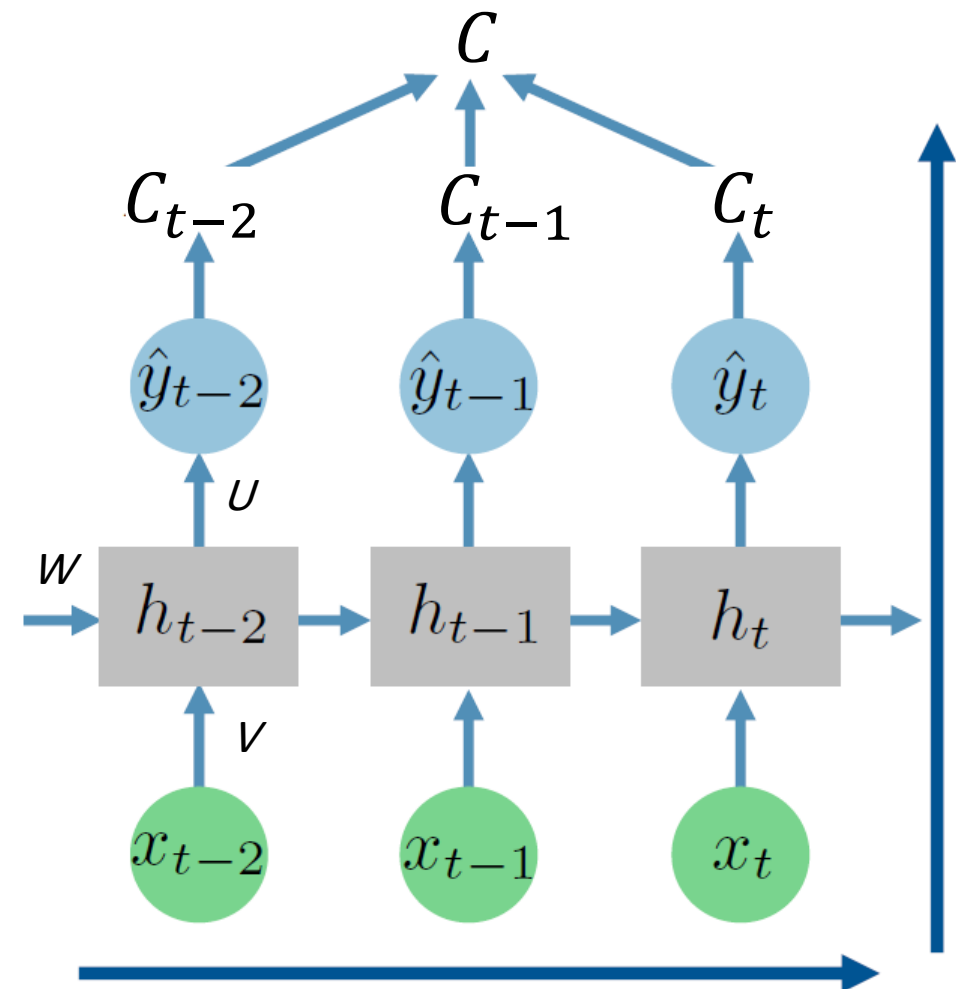
Training RNN

- ❖ Random weight initialization
- ❖ Feedforward (through time/sequence) to generate prediction, \hat{y}_t
- ❖ Compute cost function, C_t , based on actual y_t
- ❖ Total cost:

$$C = \sum_i^t C_i(y_i, \hat{y}_i)$$

- ❖ Use backprop to compute:

$$\frac{\partial C}{\partial U}, \frac{\partial C}{\partial V}, \frac{\partial C}{\partial W}, \frac{\partial C}{\partial b_y}, \frac{\partial C}{\partial b_h}$$





Backprop Through Time (BPTT)

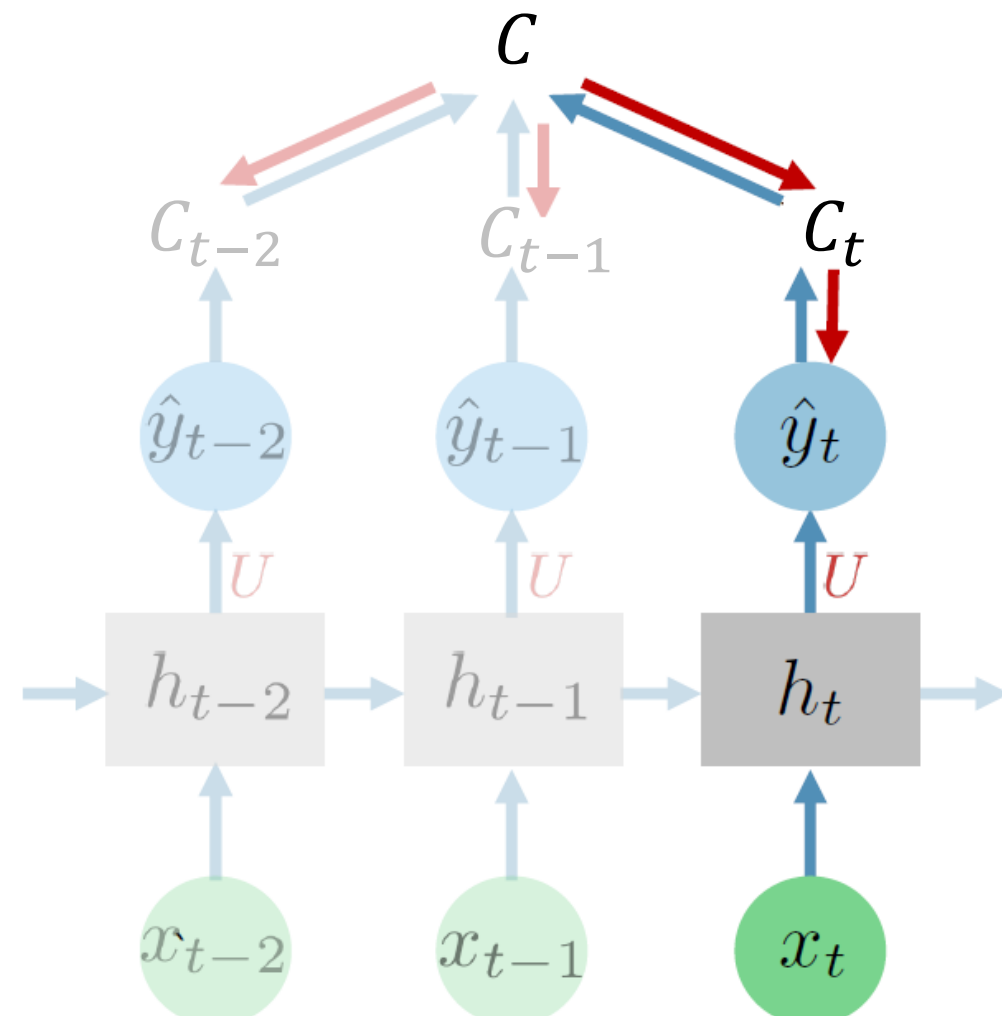
❖ Output function:

$$\hat{y}_t = f(Uh_t + b_y)$$

❖ Shared weights:

$$\frac{\partial C}{\partial U} = \sum_i^t \frac{\partial C_i}{\partial U} = \sum_i^t \frac{\partial C_i}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial U}$$

❖ For classification $\frac{\partial C_i}{\partial \hat{y}_i}$ based on the softmax





$$\frac{\partial \mathcal{C}}{\partial W}$$

$$h_t = f(Vx_t + Wh_{t-1} + b_h)$$

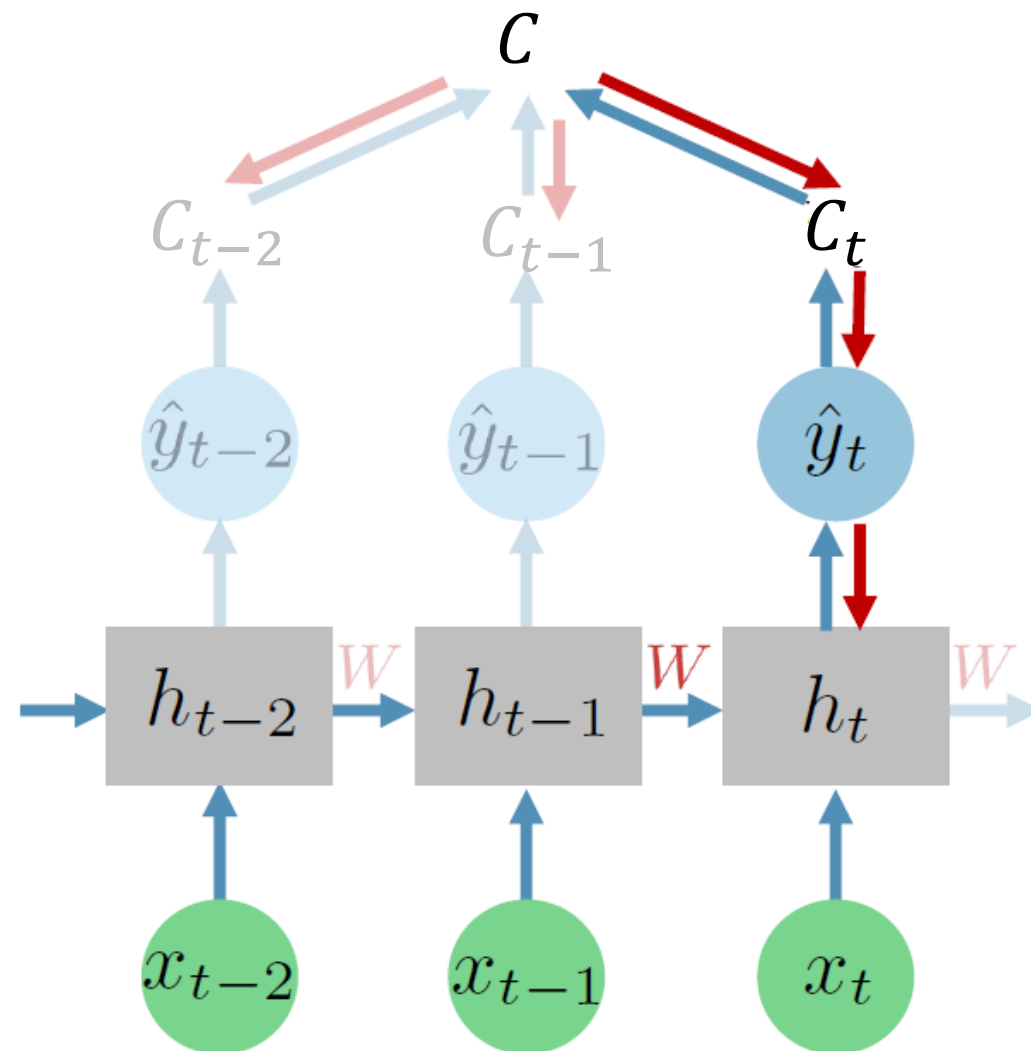
❖ h_{t-1} depends on W

$$\frac{\partial C_i}{\partial W} = \frac{\partial C_i}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial h_i} \left(\frac{\partial h_i}{\partial W} + \frac{\partial h_i}{\partial h_{i-1}} \frac{\partial h_{i-1}}{\partial W} + \dots + \frac{\partial h_1}{\partial h_0} \frac{\partial h_0}{\partial W} \right)$$

$$= \frac{\partial C_i}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial h_i} \left(\frac{\partial h_i}{\partial W} + \sum_{j=0}^{i-1} \frac{\partial h_{j+1}}{\partial h_j} \frac{\partial h_j}{\partial W} \right)$$

$$= \frac{\partial C_i}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial h_i} \left(\frac{\partial h_i}{\partial W} + \sum_{j=0}^{i-1} \left(\prod_{k=j+1}^{i-1} \frac{\partial h_{k+1}}{\partial h_k} \right) \frac{\partial h_j}{\partial W} \right)$$

$$\frac{\partial \mathcal{C}}{\partial W} = \sum_0^T \frac{\partial C_i}{\partial W}$$





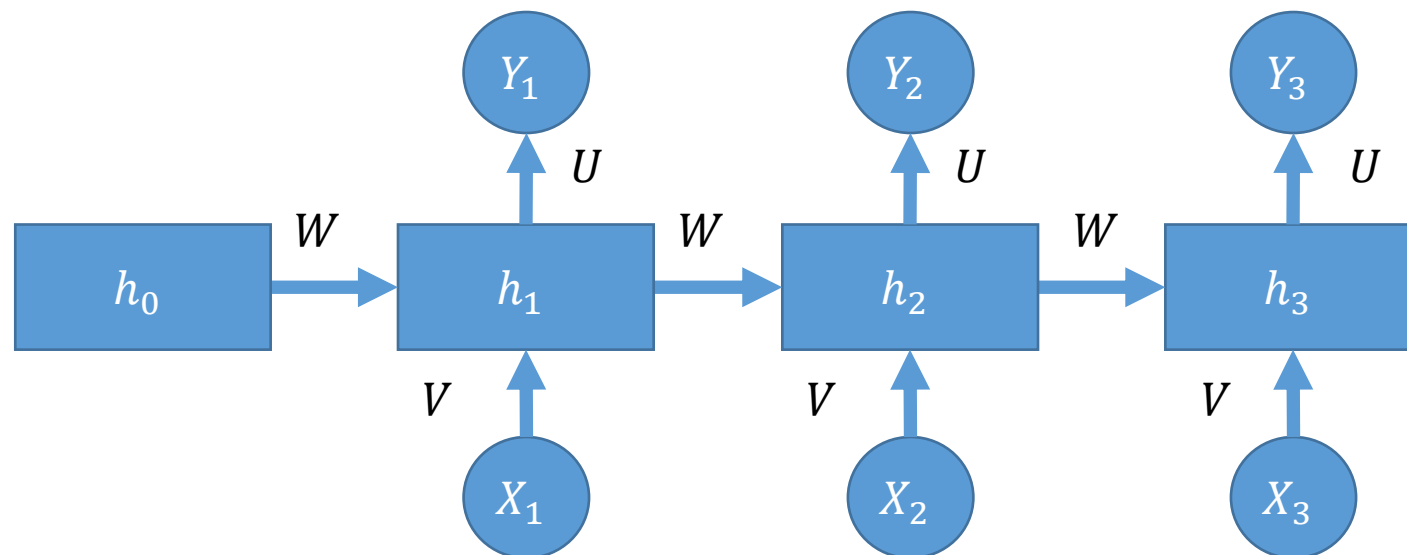
BPTT Example

$$h_1 = f(Wh_0 + Vx_1 + b_h) = f(Z_1)$$

$$h_2 = f(Wh_1 + Vx_2 + b_h) = f(Wf(Z_1) + Vx_2 + b_h) = f(Z_2)$$

$$= (W^2h_0 + WVx_1 + Wb_h) + Vx_2 + b_h$$

$$h_3 = f(Wh_2 + Vx_3 + b_h) = f(Wf(Z_2) + Vx_3 + b_h) = f(Z_3)$$





$$h_n$$

$$h_1 = f(Wh_0 + Vx_1 + b_h)$$

$$h_2 = f(Wh_1 + Vx_2 + b_h) = (W^2h_0 + WVx_1 + Wb_h) + Vx_2 + b_h$$

$$\begin{aligned} h_3 &= Wh_2 + Vx_3 + b_h = W((W^2h_0 + WVx_1 + Wb_h) + Vx_2 + b_h) + Vx_3 + b_h \\ &= W^3h_0 + W^2Vx_1 + W^2b_h + WVx_2 + Wb_h + Vx_3 + b_h \\ &= W^3h_0 + W^2Vx_1 + WVx_2 + Vx_3 + W^2b_h + Wb_h + b_h \end{aligned}$$

Let $x_0 = h_0$:

$$h_3 = W^3x_0 + W^2Vx_1 + WVx_2 + Vx_3 + W^2b_h + Wb_h + b_h$$

$$h_n = W^n x_0 + \sum_{t=1}^n W^{n-t} Vx_t + W^{n-t} b_h$$



$$\frac{\partial h_3}{\partial W}$$

$$\frac{\partial h_3}{\partial W} = \frac{\partial f(Z_3)}{\partial Z_3} \frac{\partial Z_3}{\partial g(Z_2)} \frac{\partial f(Z_2)}{\partial Z_2} \frac{\partial Z_2}{\partial g(Z_1)} \frac{\partial f(Z_1)}{\partial Z_1} \frac{\partial Z_1}{\partial h_0}$$

$$\frac{\partial h_3}{\partial W} = \frac{\partial h_3}{\partial Z_3} \frac{\partial Z_3}{\partial h_2} \frac{\partial h_2}{\partial Z_2} \frac{\partial Z_2}{\partial h_1} \frac{\partial h_1}{\partial Z_1} \frac{\partial Z_1}{\partial h_0} = f'(Z_3)W f'(Z_2)W f'(Z_1)W$$

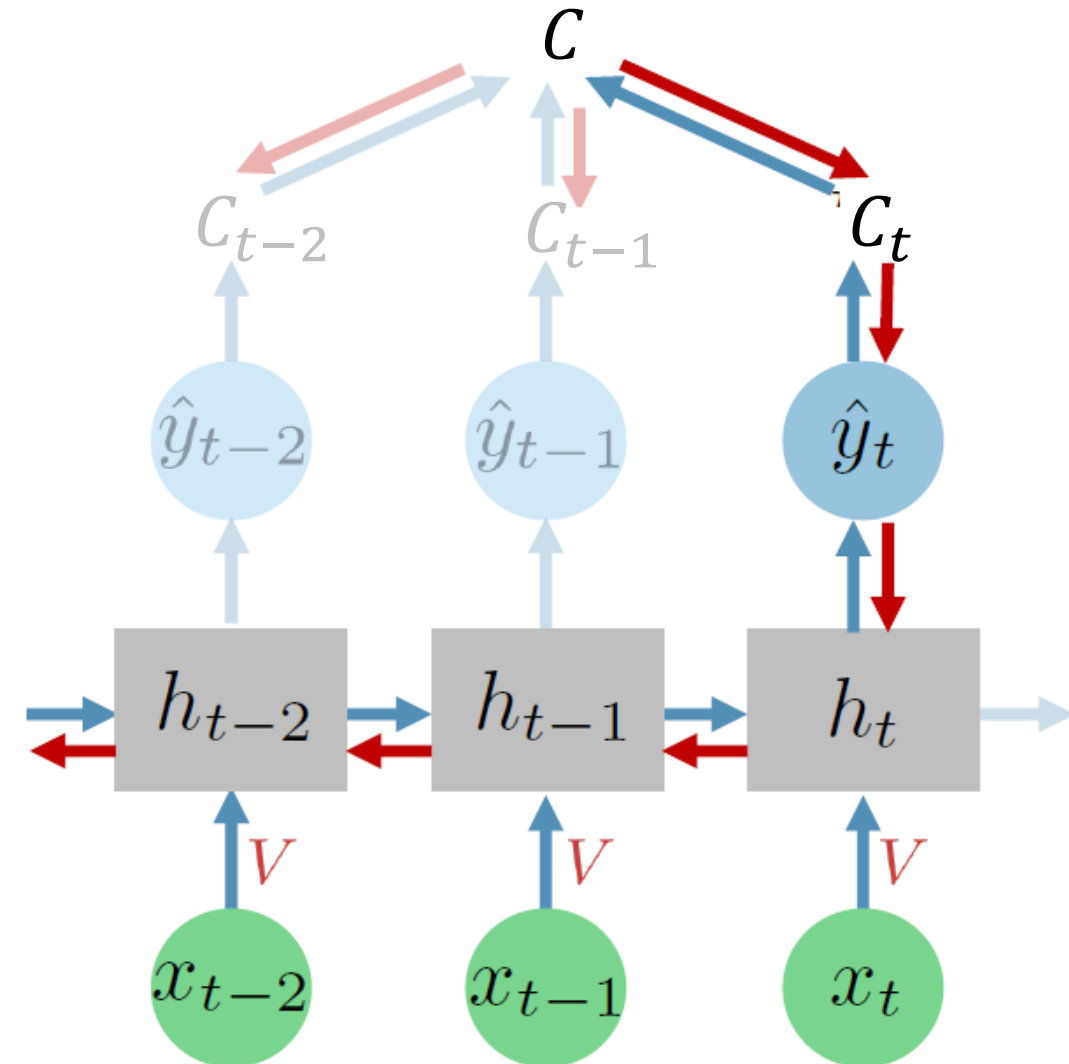


$$\frac{\partial \mathcal{C}}{\partial V}$$

$$h_t = f(Vx_t + Wh_{t-1} + b_h)$$

❖ Similarly, h is dependent on V

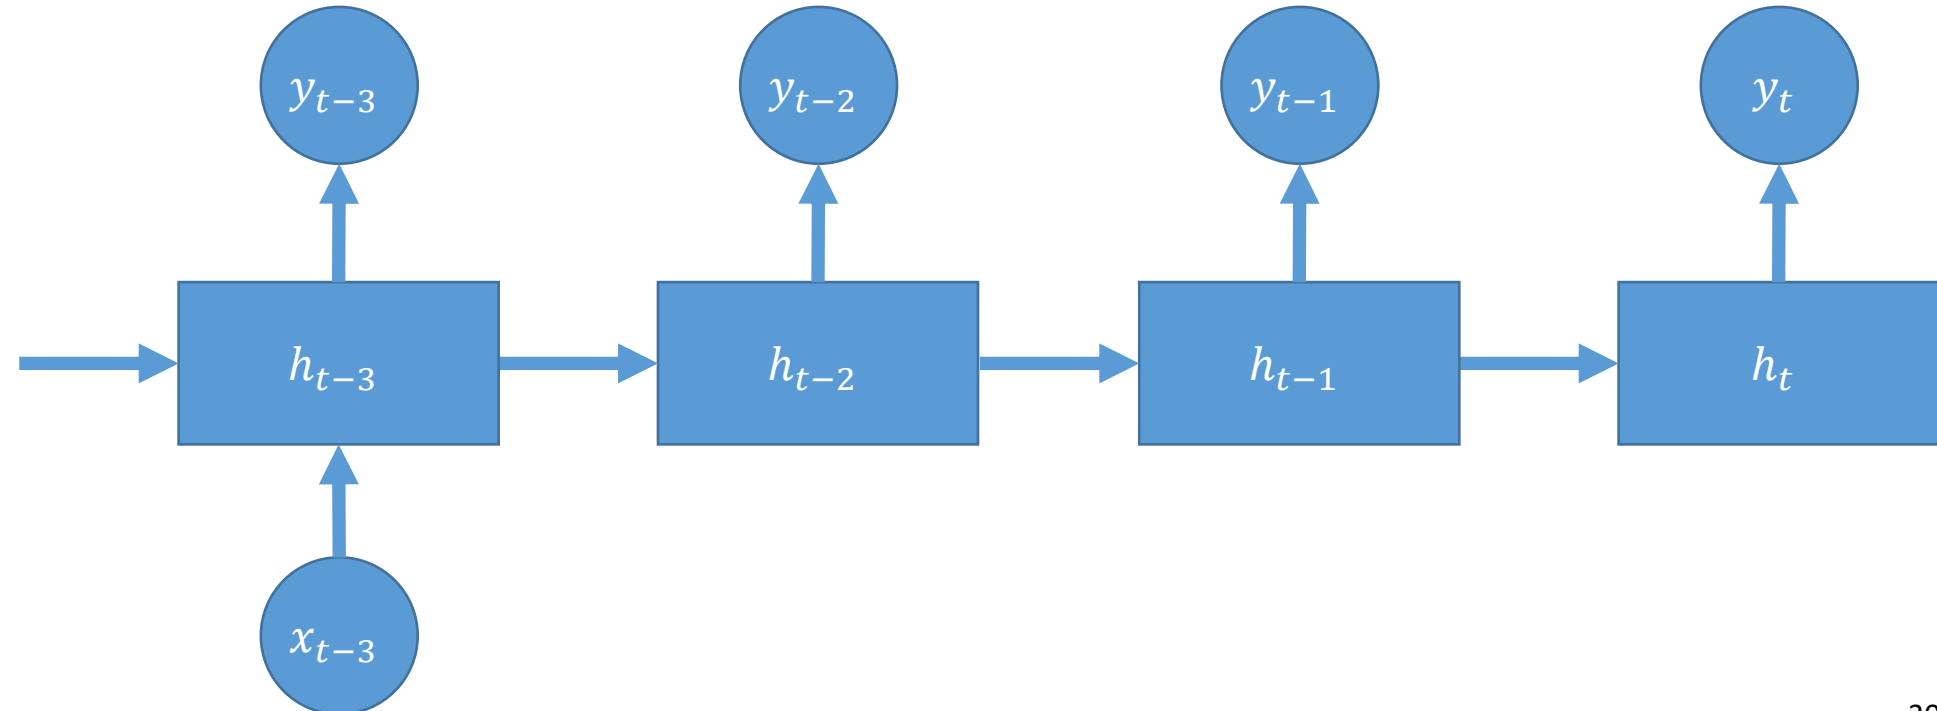
$$\frac{\partial \mathcal{C}_i}{\partial V} = \frac{\partial \mathcal{C}_i}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial h_i} \left(\frac{\partial h_i}{\partial V} + \sum_{j=0}^{i-1} \left(\prod_{k=j+1}^{i-1} \frac{\partial h_{k+1}}{\partial h_k} \right) \frac{\partial h_j}{\partial V} \right)$$





RNN One input – Many Outputs

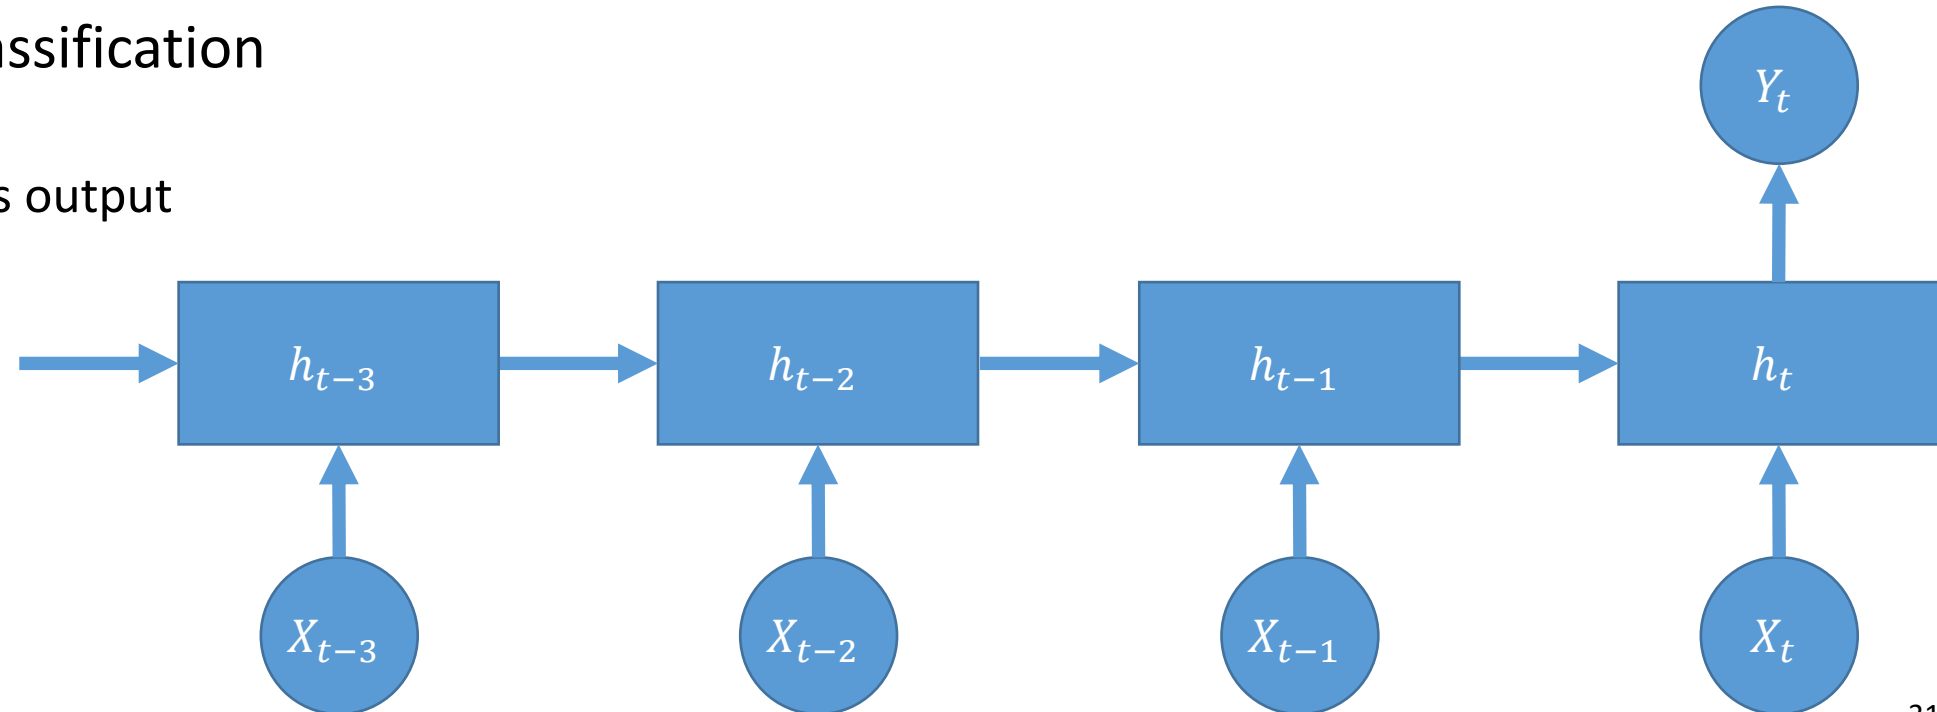
- ❖ Sequence generation
- ❖ Example: image caption
 - Image as input
 - Text as output





RNN Many inputs – One output

- ❖ Sequence based classification/prediction
- ❖ Example: speech recognition
 - Audio clips as input
 - Word as output
- ❖ Example: text classification
 - Words as input
 - Subject/topic as output





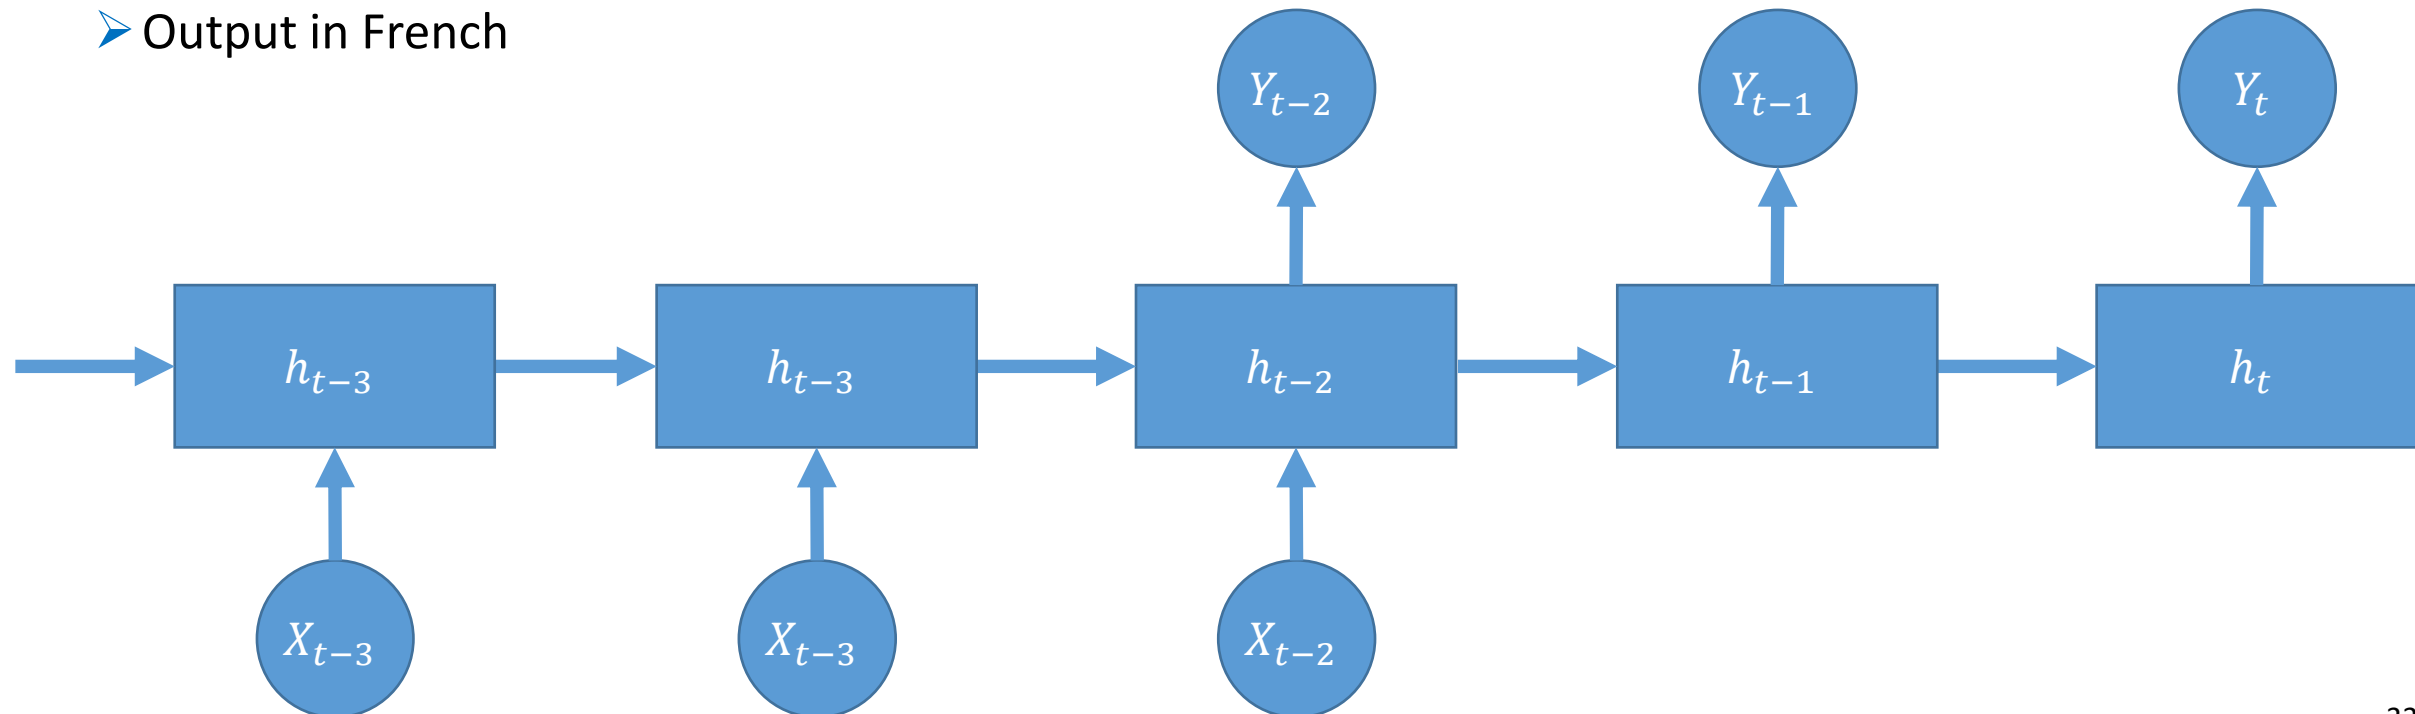
RNN Many inputs – Many outputs

❖ Delayed output

- Encoder-Decoder design

❖ Example: machine translation

- Input in English
- Output in French



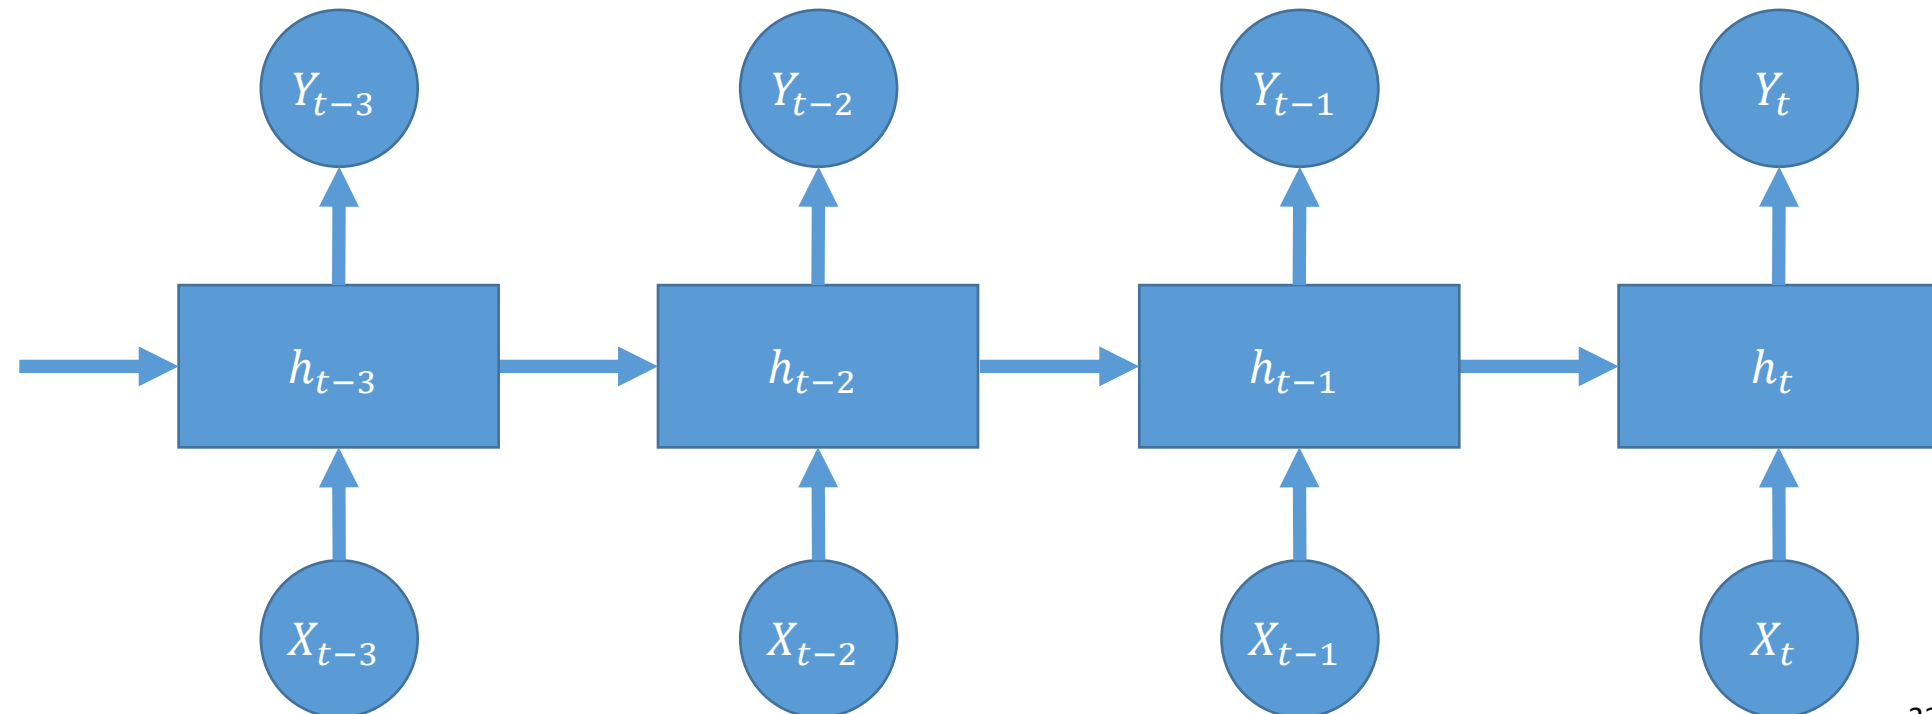


RNN Many inputs – Many outputs

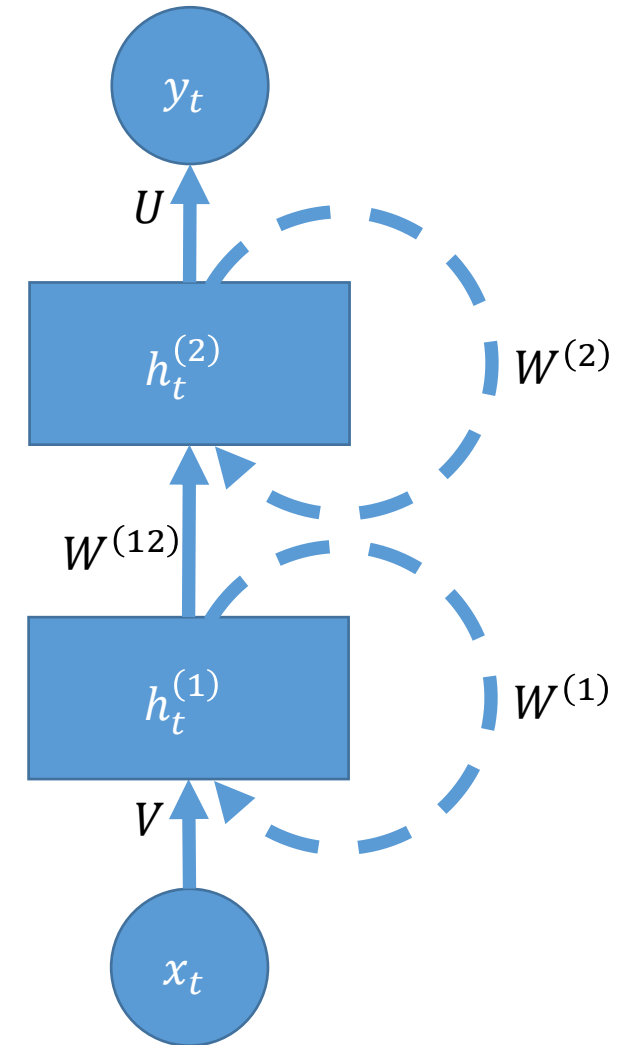
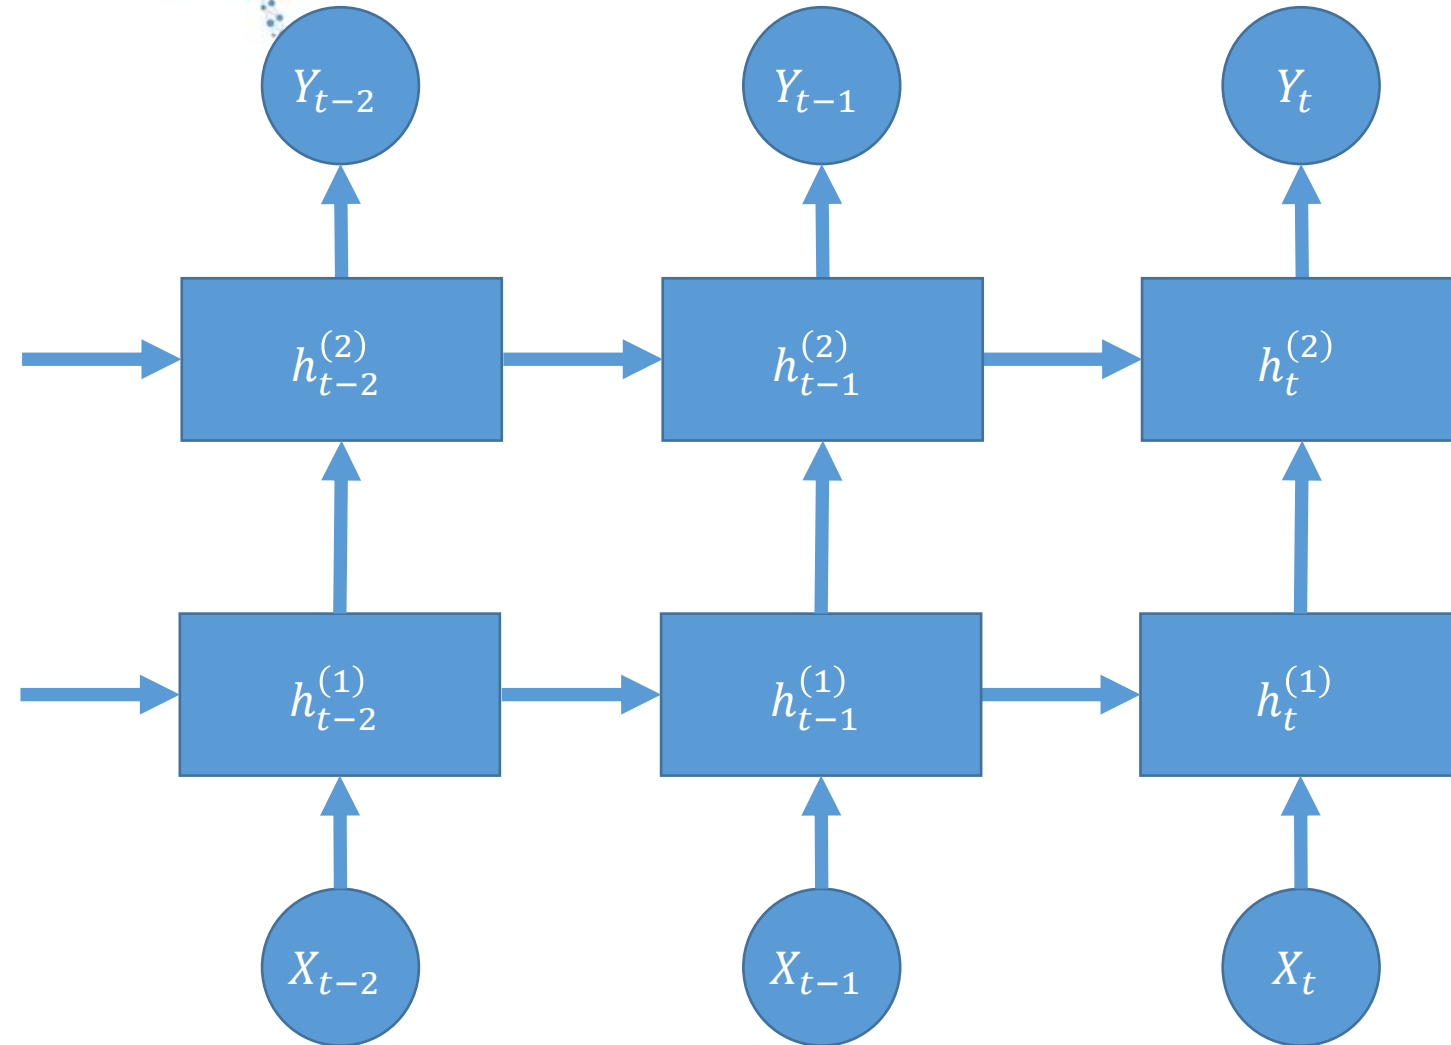
❖ Synched output

❖ Example: stock prediction

- Value of stock is fed in at each iteration
- Predicted value of stock is outputted at each sequence

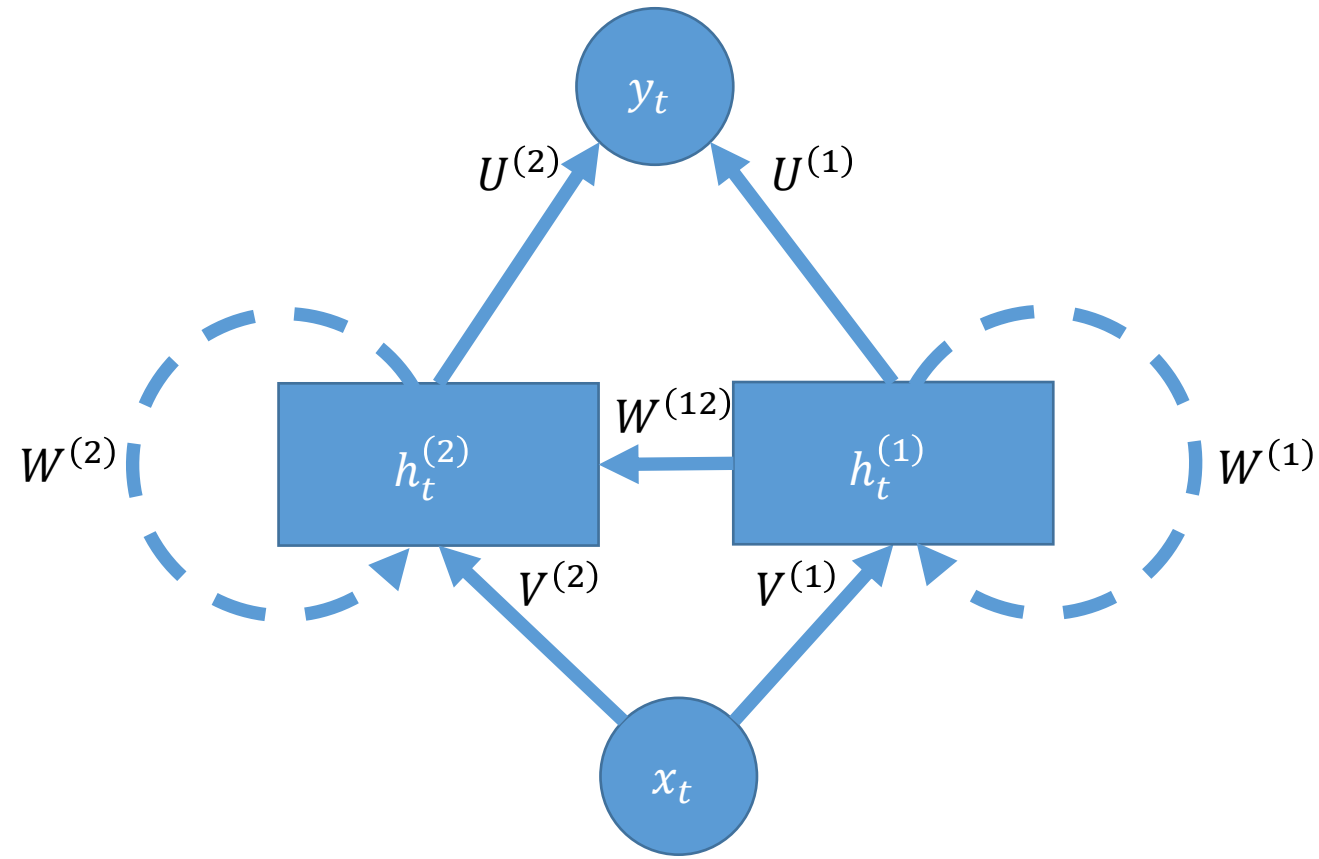


Sequential Recursion



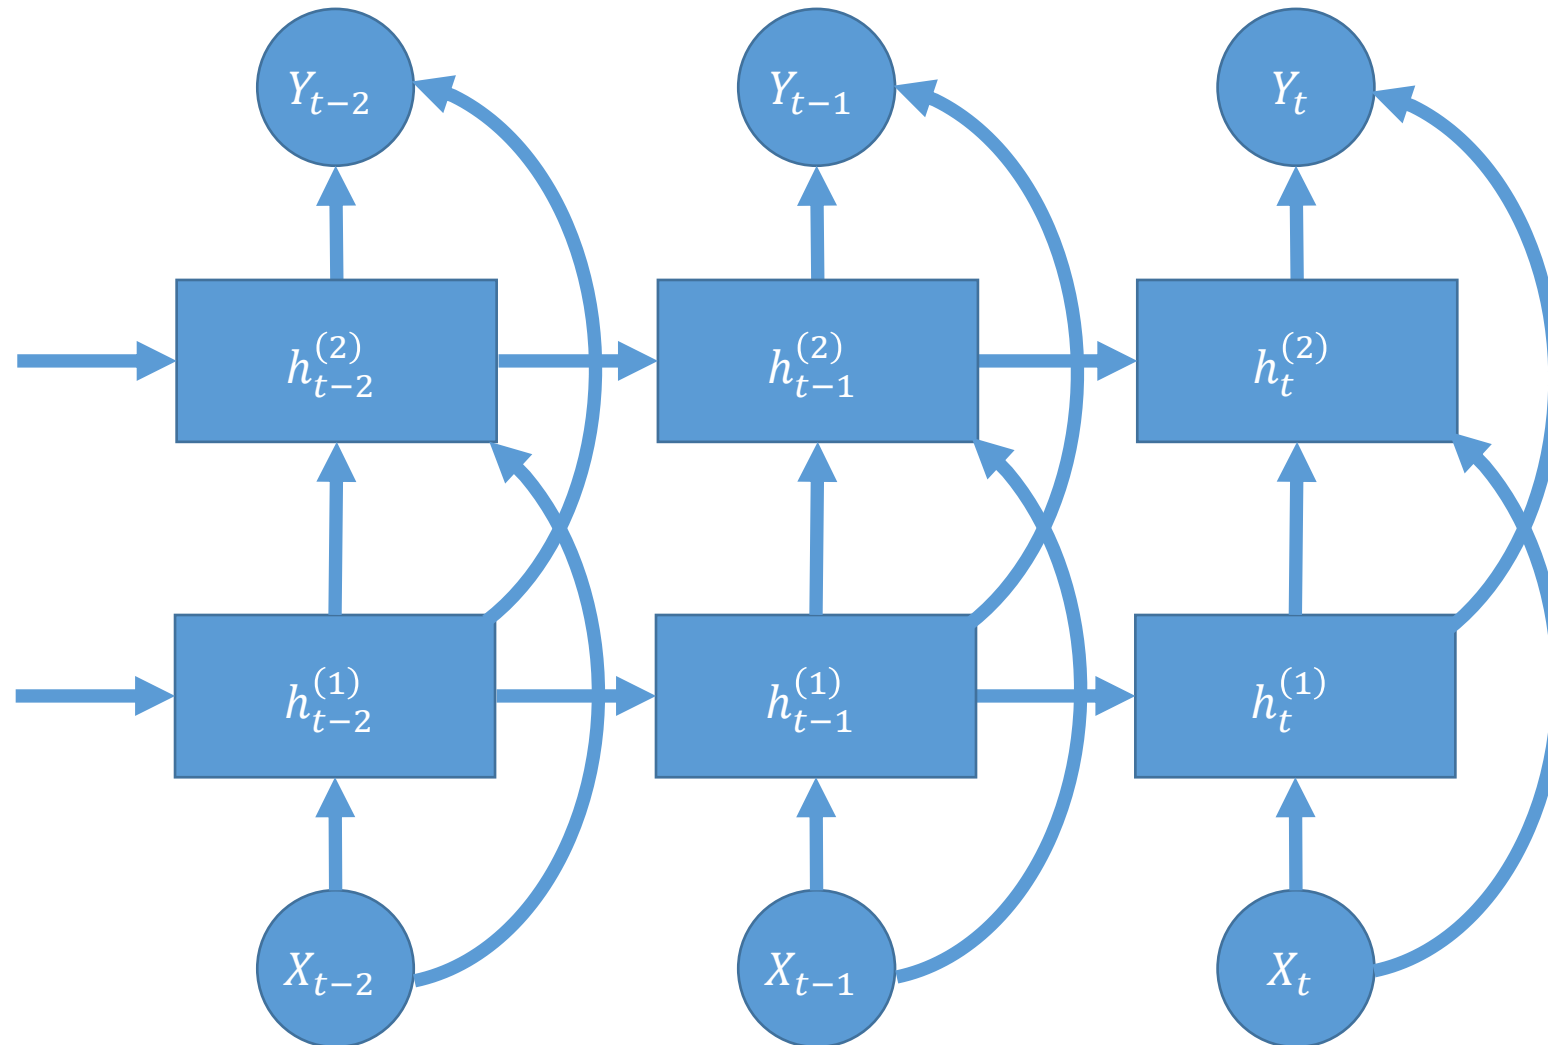


Parallel Recursion





Unfolded Parallel Recursion





BRNN (1997)

❖ Bi-Directional RNN: Parallel backward and forward recursions

- Schuster and Paliwal
- especially useful when the context of the input

❖ Forward layer

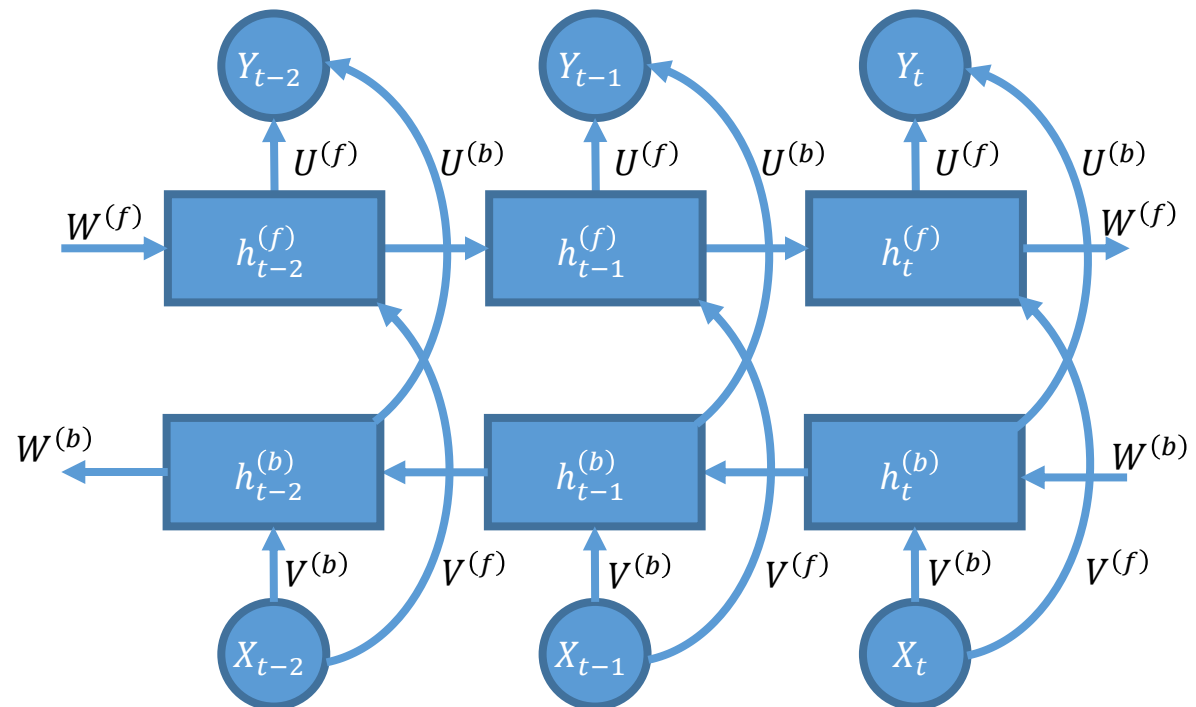
- Predict future from past
- Processes data from $t=0$ to T

❖ Backward layer

- Deduces past from future
- Processes data from $t=T$ to 0

❖ Applications:

- Speech Recognition
- Translation
- Handwritten Recognition
- Protein Structure Prediction





Stability Analysis

- ❖ Problem: recursion and output saturation
 - For sigmoid/Tanh: saturation
 - For ReLU: explosion
- ❖ Output stability, single tap: $y_t = g(Uh_t)$
 - if h_t is bounded (stable) then output is stable
 - Ignoring bias
- ❖ Hidden layer stability, single tap: $h_t = f(Vx_t + Wh_{t-1})$
 - x_t bounded (naturally)
 - Process depends on recursion

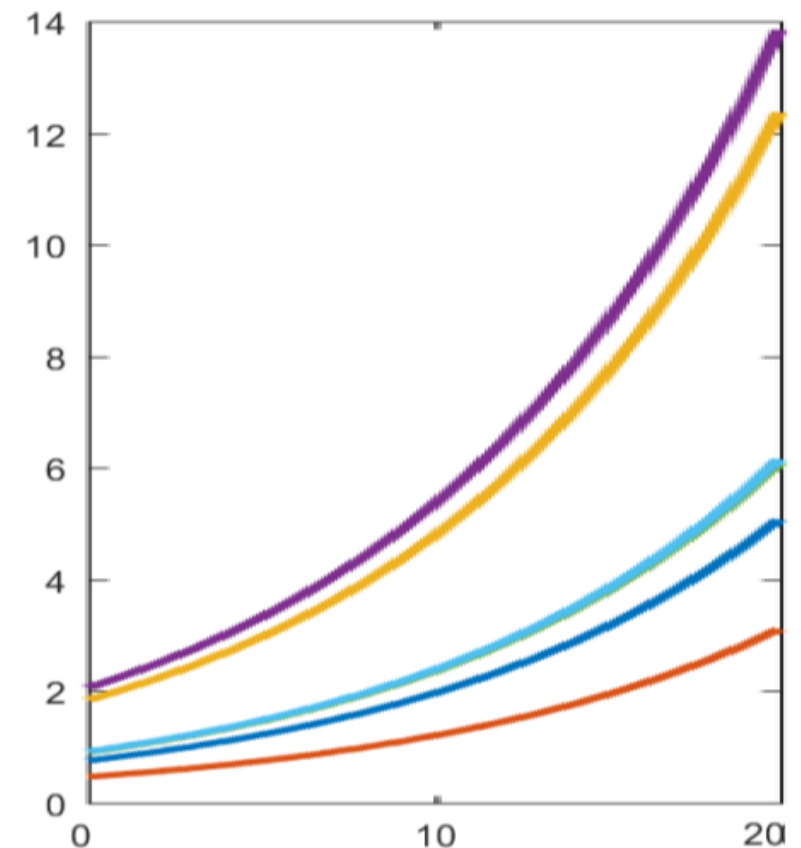
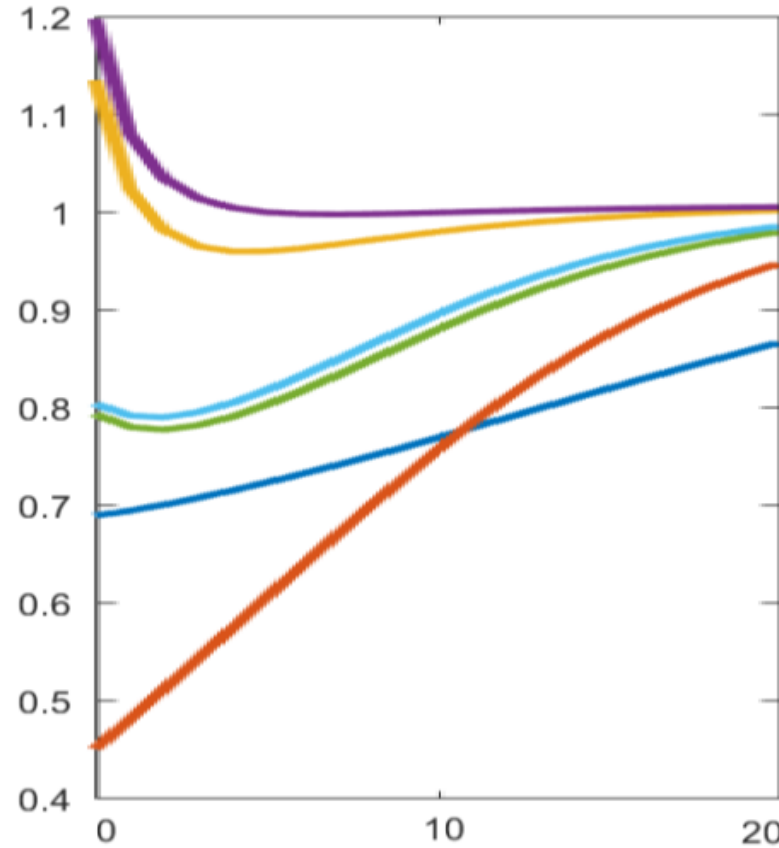
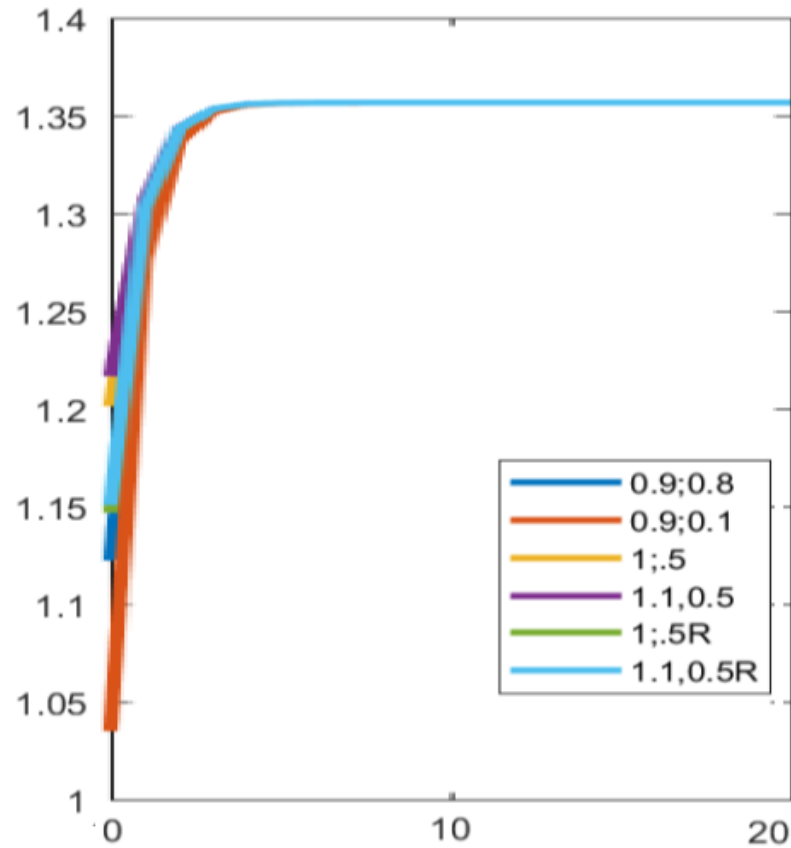


Sigmoid

Stability of AF for 1 initialization

Tanh

ReLU





Stability and Memory

- ❖ Weight of recursion can cause instability
- ❖ Bipolar functions hold memory
- ❖ Low stability \Rightarrow memory is low
- ❖ Exponential instability \Rightarrow memory is forgotten exponentially



Gradient Stability

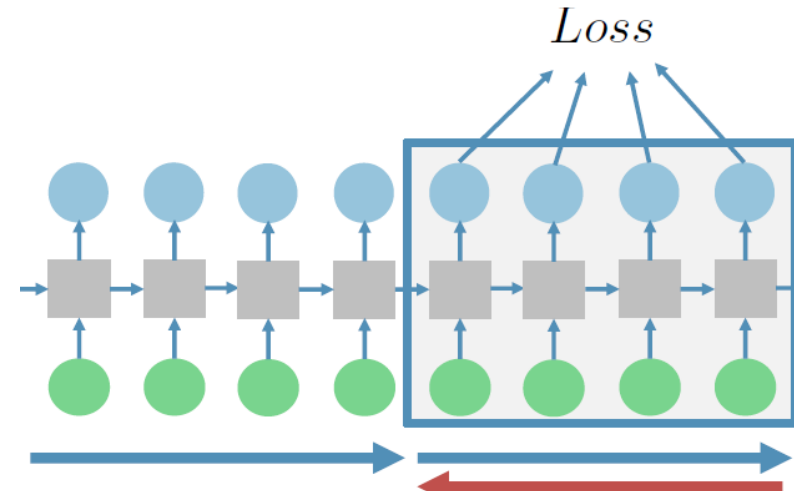
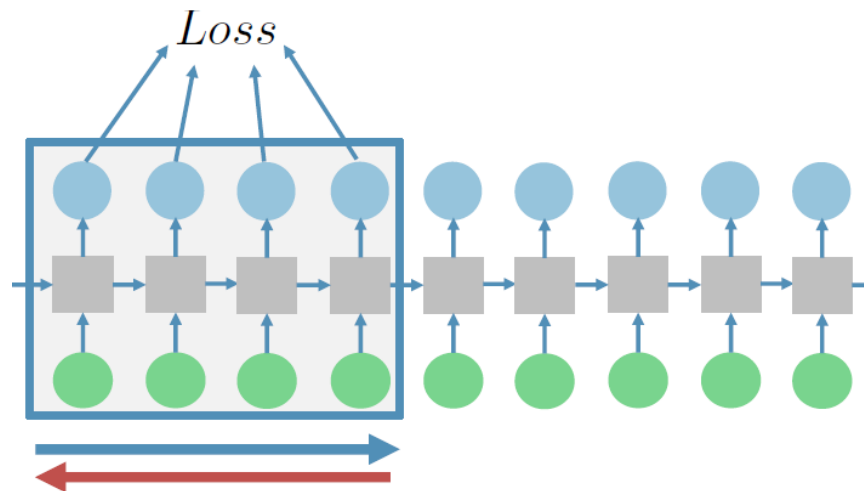
❖ Deep recursions synonymous with deep networks

➤ Gradients will explode or vanish

❖ Exploding Solutions:

➤ Clipping grads: if $\|g\| > Thresh \Rightarrow g = \frac{Thresh}{\|g\|} g$

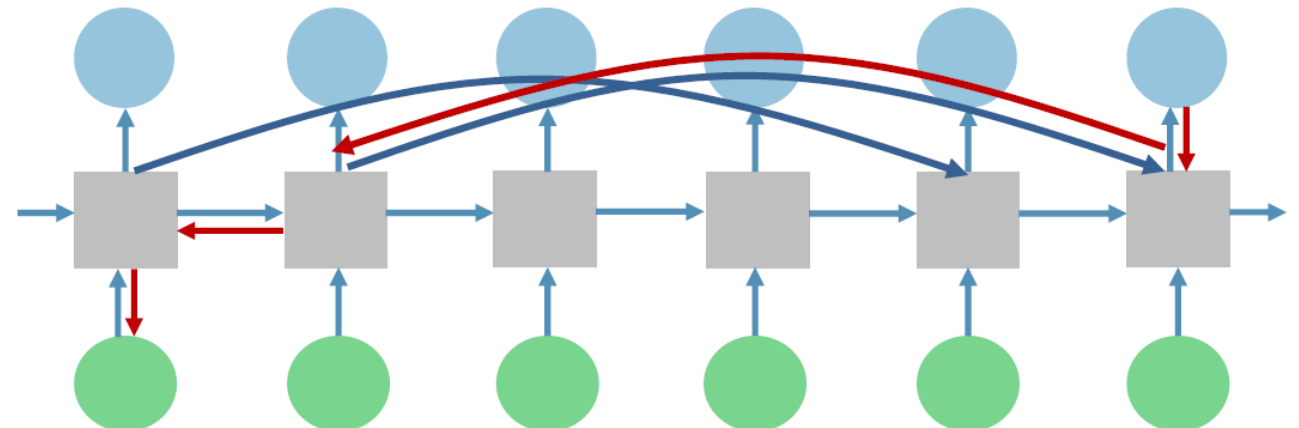
➤ Batch Process Loss: Forward pass and backward pass chunks





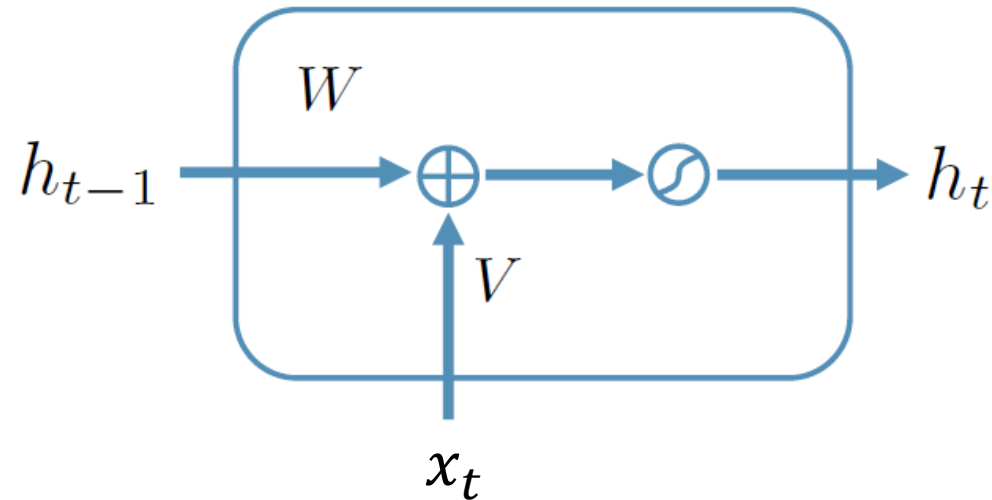
Vanishing Grad Solutions

- ❖ Vanishing grads \Rightarrow loss of memory
- ❖ ReLU activation
 - Sigmoid/tanh: saturate and cause gradient to vanish
- ❖ Initialization solutions
 - Choose orthogonal W matrix: $W^T = W^{-1}$
 - $w_{ij}^{-1} = w_{ji}$
 - Orthogonal W doesn't vanish (or explode)
- ❖ Skip connections
 - BP across skip connections
 - Vanishes slower than other connections





Hidden State as a Cell



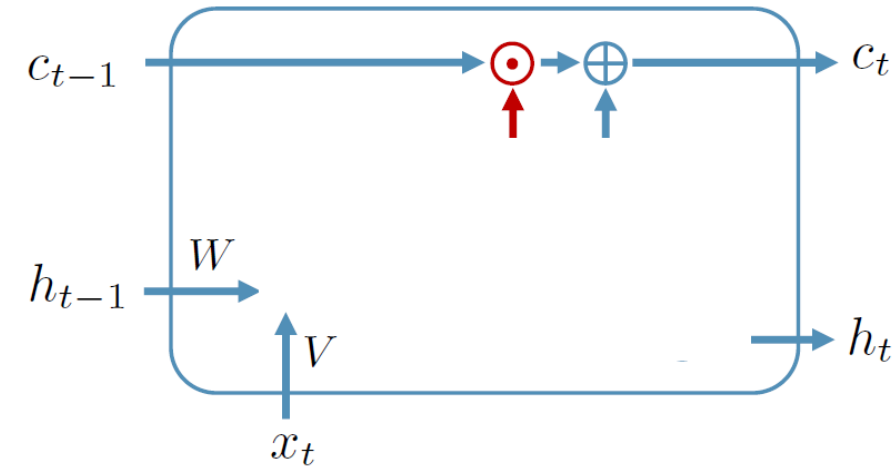
$$h_t = f(W h_{t-1} + V x_t + b_h)$$

- ❖ Nonlinearity causes vanishing gradient in the backward-pass
- ❖ Memory and gradient are tied together
 - Vanishing memory \Rightarrow lost memory



LSTM: Memory as a Separate Path

- ❖ Long Short Term Memory
- ❖ Idea1 : Separate path for cell memory
 - Allows easier derivative ($\partial c_t / \partial c_{t-1} = 1$)
 - No nonlinearities in the path
- ❖ c_t is cell memory
 - Same dimension as h_t
- ❖ Idea 2: Update information from the cell into memory
 - Input & previous states contains information
 - Updates can be added \oplus
 - Updates can be scaled \odot



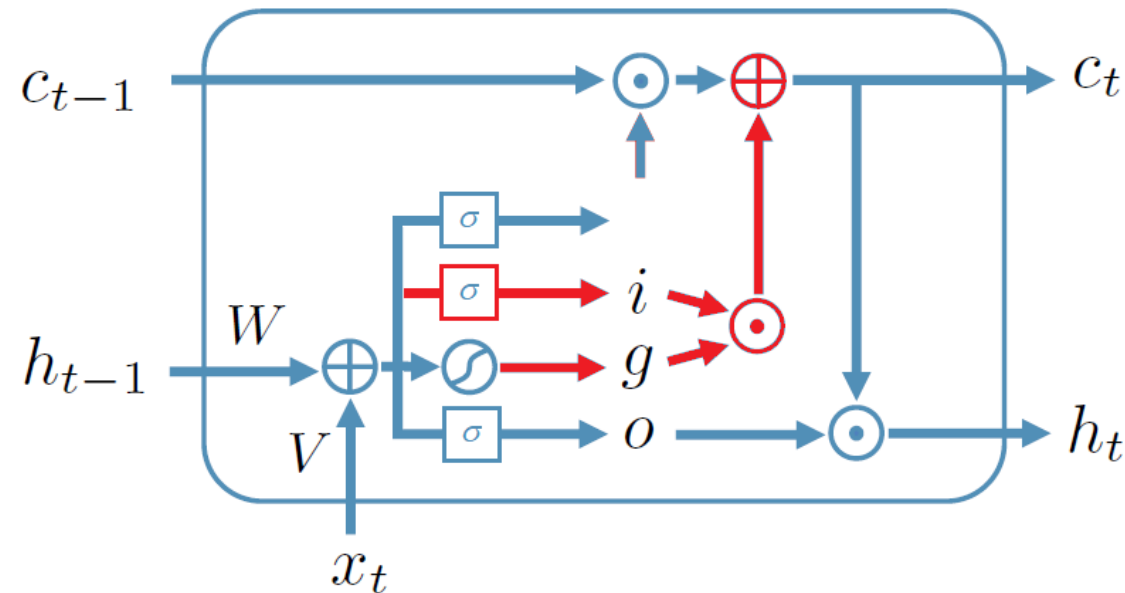


LSTM: Memory Updates

- ❖ Idea 2: Update memory with information from the cell into memory
 - Based on input and previous state
- ❖ Updates can be added \oplus
 - Percentage of x_t & h_{t-1}
- ❖ Updates can be multiplicative \odot
 - Cause cell to forget
 - Use sigmoid function, $\sigma(\quad)$
 - Need a percentage for scaling purposes



LSTM: Update Memory with new Input



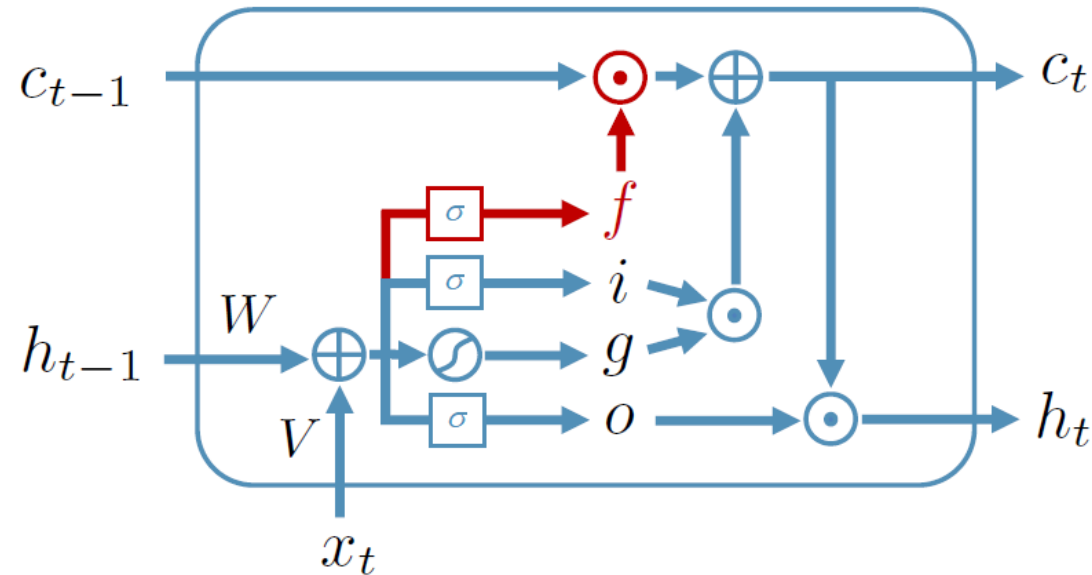
$$\begin{aligned}
 i &= \sigma(W h_{t-1} + V x_t + b_i) \\
 g &= \tanh(W h_{t-1} + V x_t + b_g) \\
 c_t &= i \odot g + c_{t-1}
 \end{aligned}$$

❖ Input: $i \in [0: 1]$

➤ Controls how much of past memory moves forward



LSTM: Forget Past Memory



$$f = \sigma(W h_{t-1} + V x_t + b_f)$$

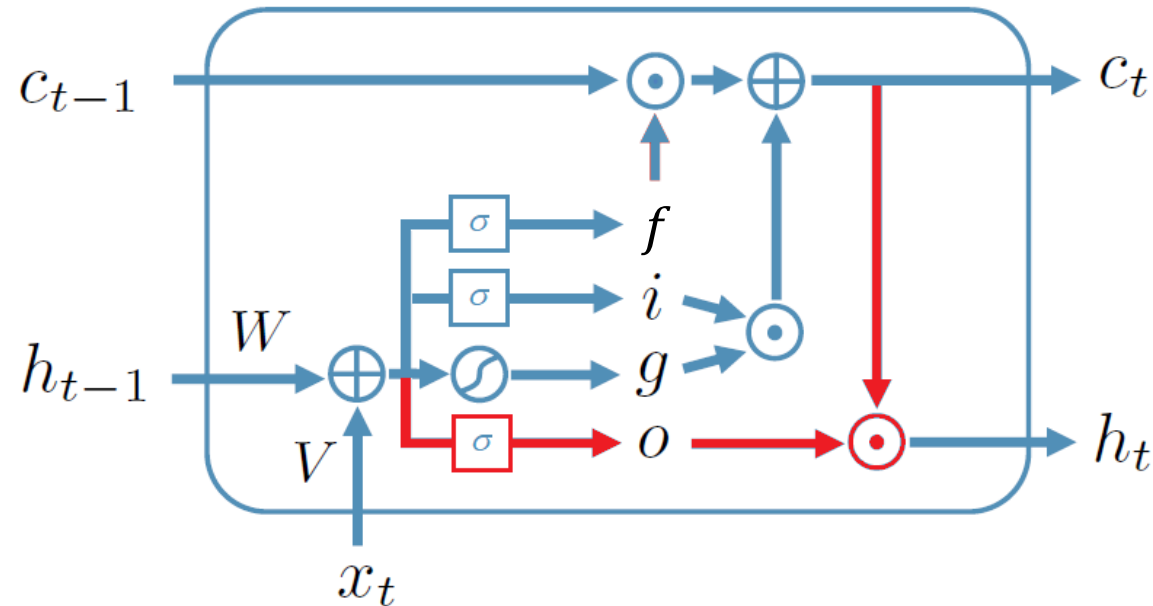
$$c_t = f \odot c_{t-1} + i \odot g$$

❖ Forget: $f \in [0: 1]$

➤ Controls how much of past memory moves forward



LSTM: Output



$$o = \sigma(W h_{t-1} + V x_t + b_o)$$

$$h_t = o \odot c_t$$

❖ Output: $o \in [0: 1]$

➤ Controls how much of new memory is encoded in new state



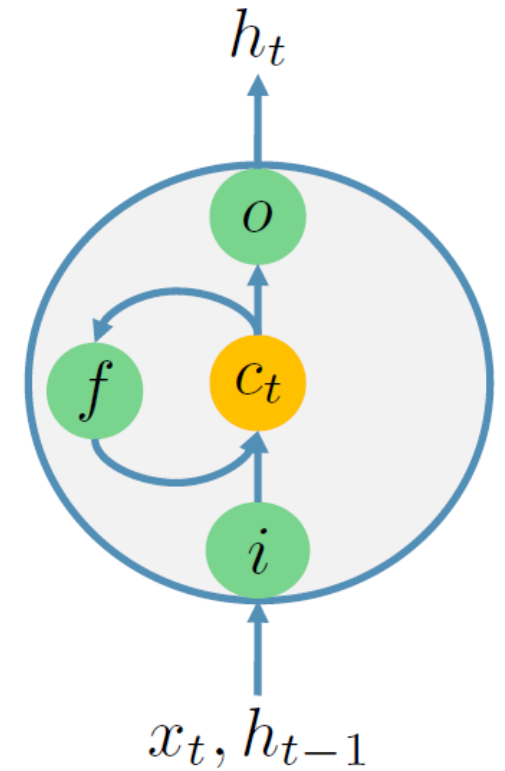
LSTM Cell Model

$$i = \sigma(W h_{t-1} + V x_t + b_i)$$

$$c_t = i \odot g + f \odot c_{t-1}$$

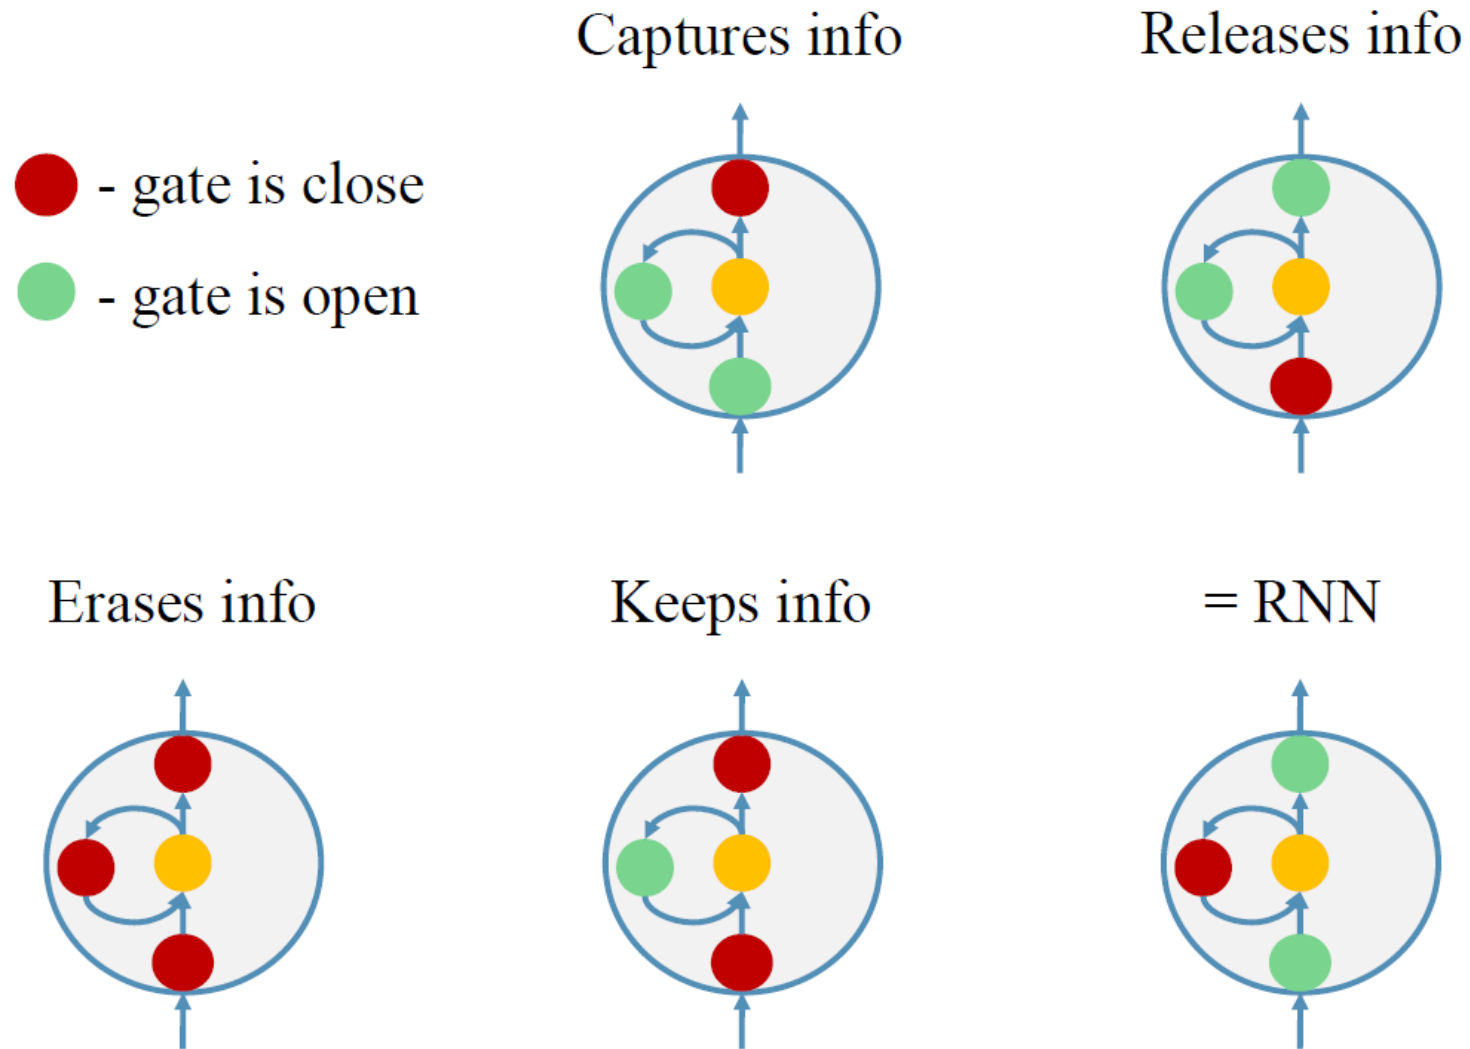
$$f = \sigma(W h_{t-1} + V x_t + b_f)$$

$$o = \sigma(W h_{t-1} + V x_t + b_o) \odot c_t$$



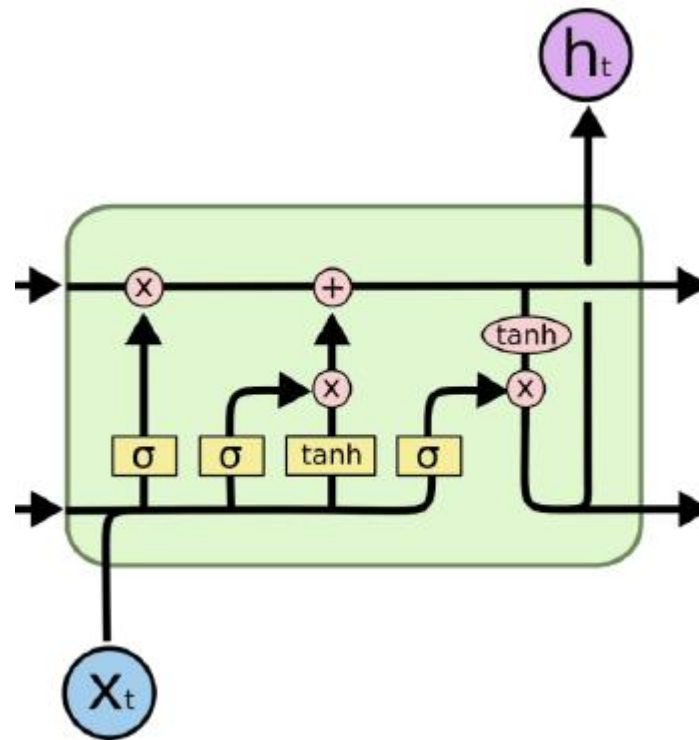


Extreme Conditions





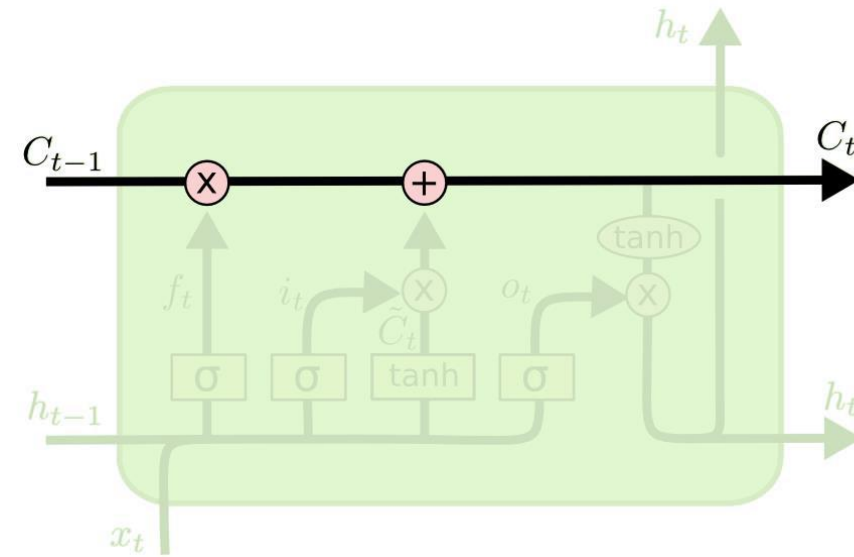
Alternative LSTM2



<http://colah.github.io/posts/2015>

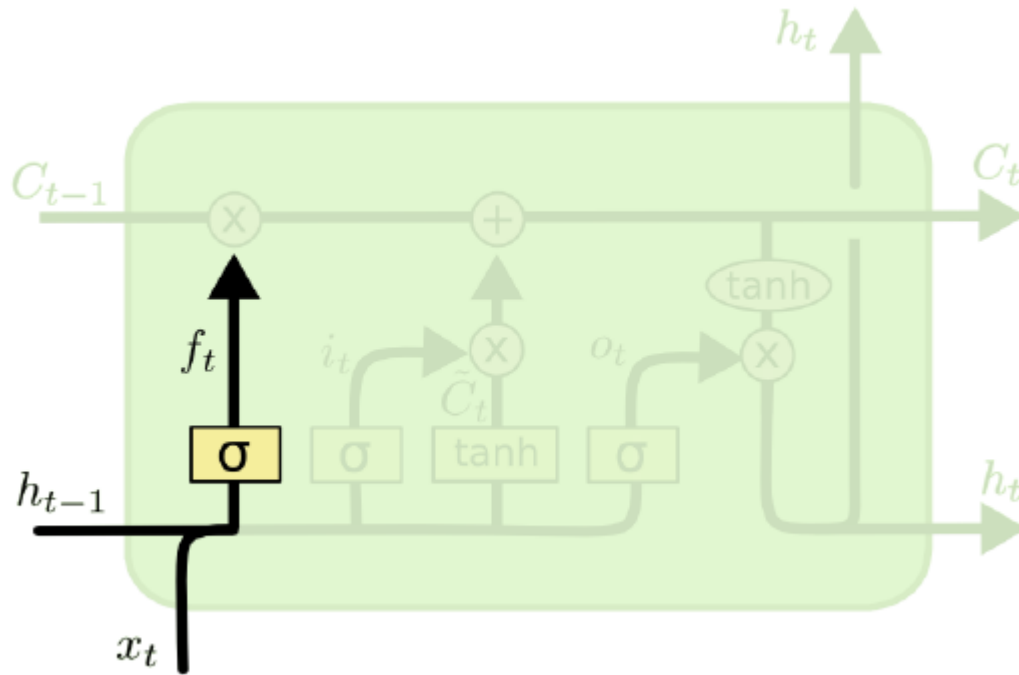


LSTM2: Memory Path





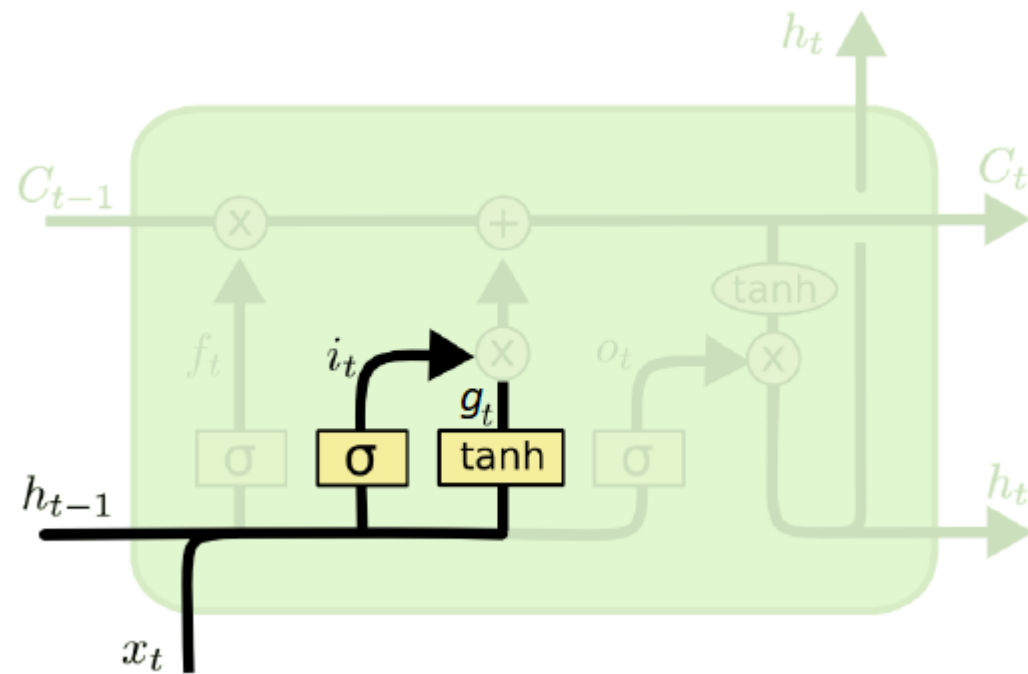
LSTM2: Forget Gate



$$f_t = \sigma(W h_{t-1} + V x_t + b_f)$$



LSTM2: Input Gate

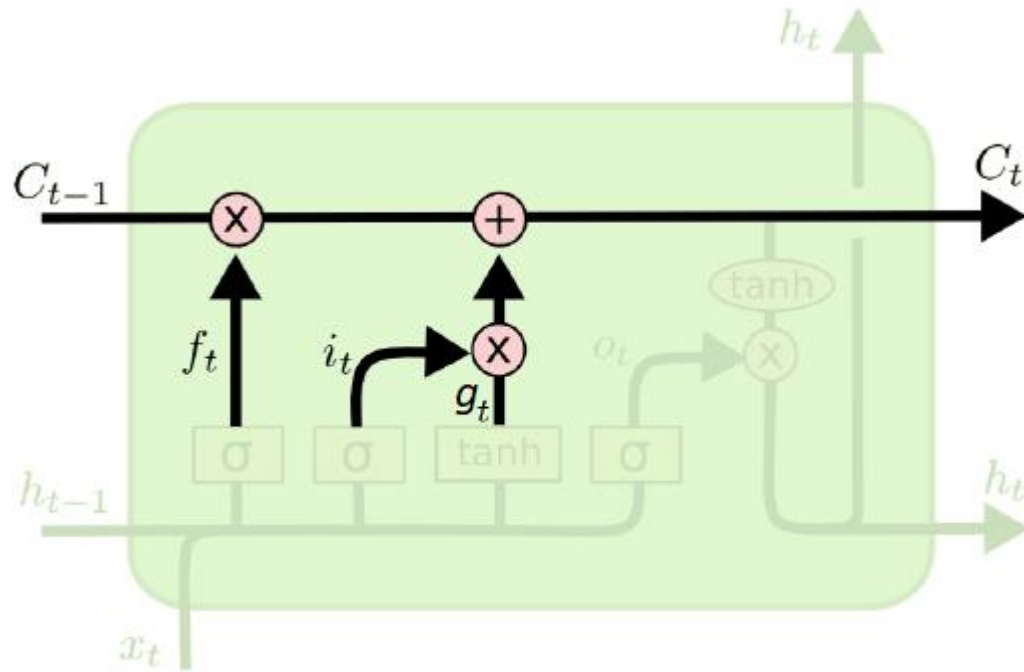


$$i_t = \sigma(W_i h_{t-1} + V_i x_t + b_i)$$

$$g_t = \tanh(W_g h_{t-1} + V_g x_t + b_g)$$



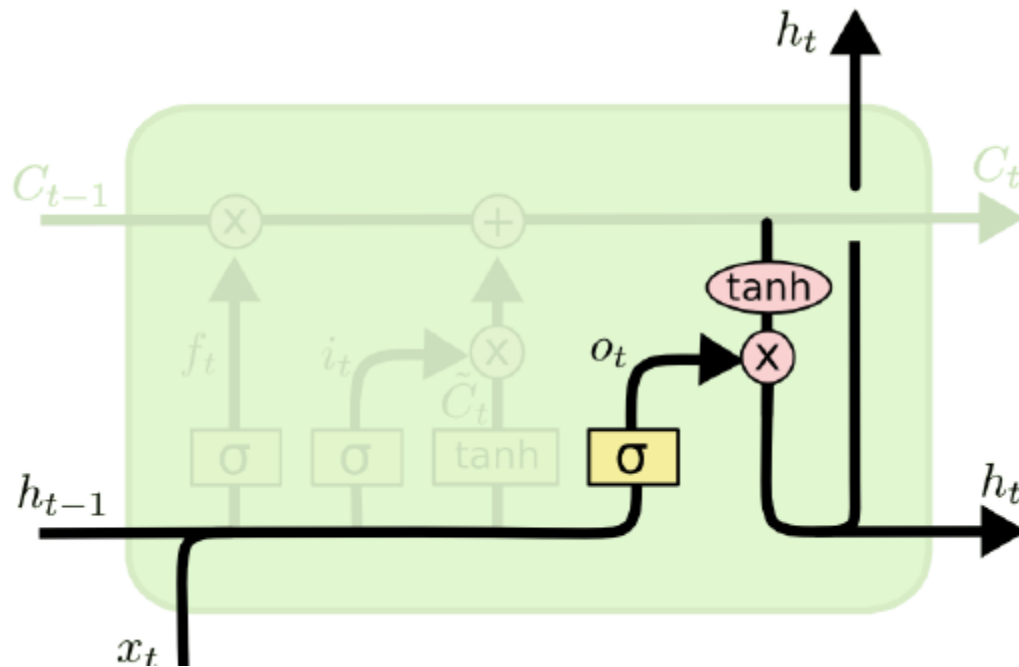
LSTM2: Update



$$c_t = f_t c_{t-1} + i_t g_t$$



LSTM2: Output Gate



$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$



LSTM2: Cell Model

$$i_t = \sigma(W_i h_{t-1} + V_i x_t + b_i)$$

$$f_t = \sigma(W_f h_{t-1} + V_f x_t + b_f)$$

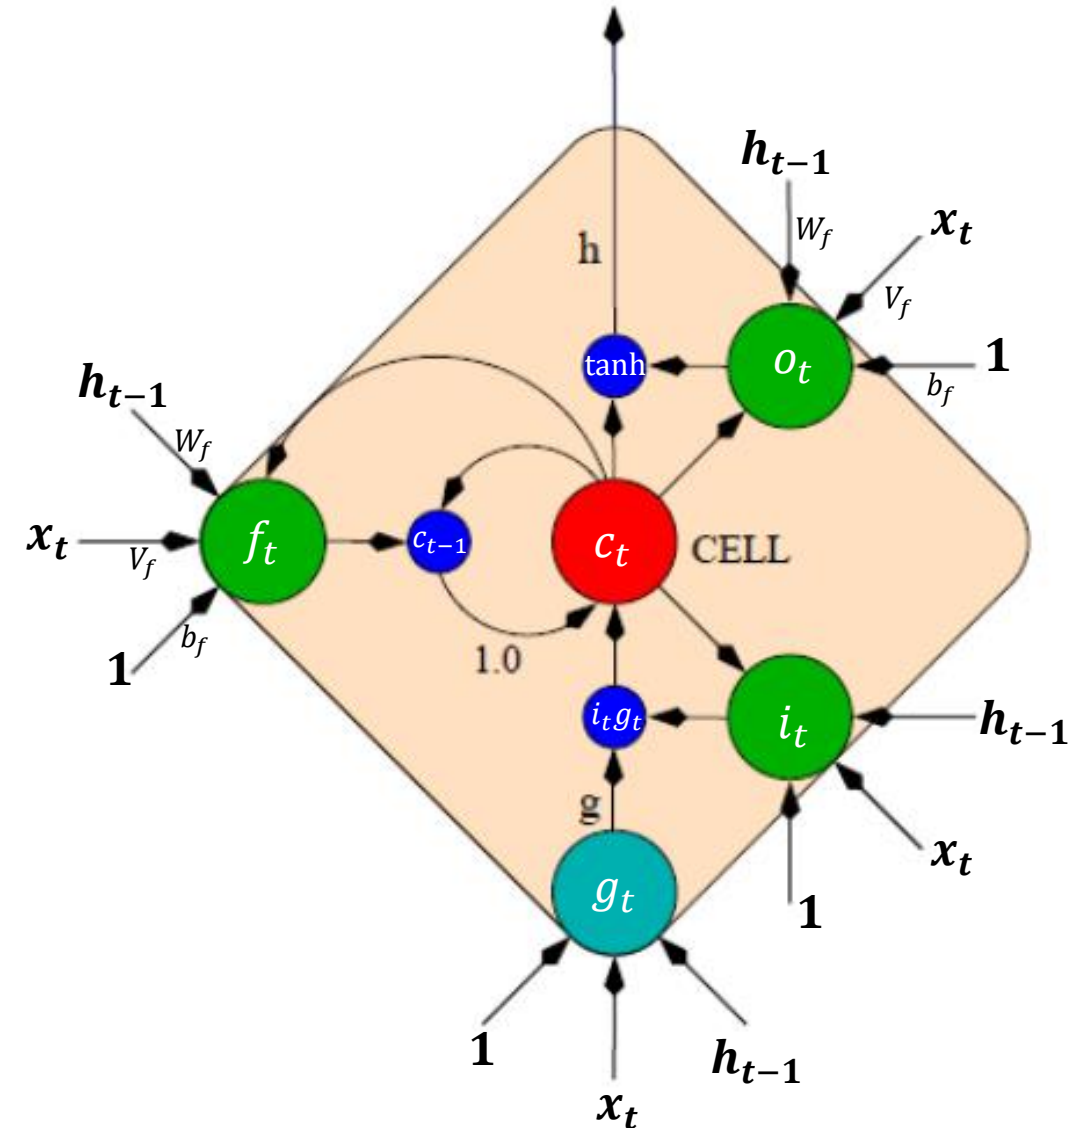
$$o_t = \sigma(W_h h_{t-1} + V_h x_t + b_h)$$

$$g_t = \tanh(W_g h_{t-1} + V_g x_t + b_g)$$

$$c_t = f_t c_{t-1} + i_t g_t$$

$$h_t = o_t \tanh(c_t)$$

$$y_t = \text{softmax}(U h_{t-1} + b_y)$$





❖ LSTM

- Outputs: c, h
- States: input, output, forget
- Additional parameters: g, i, o, f

❖ Gated Recurrent Units

- Output: h
- States: output, forget
- Additional Parameters: r, u

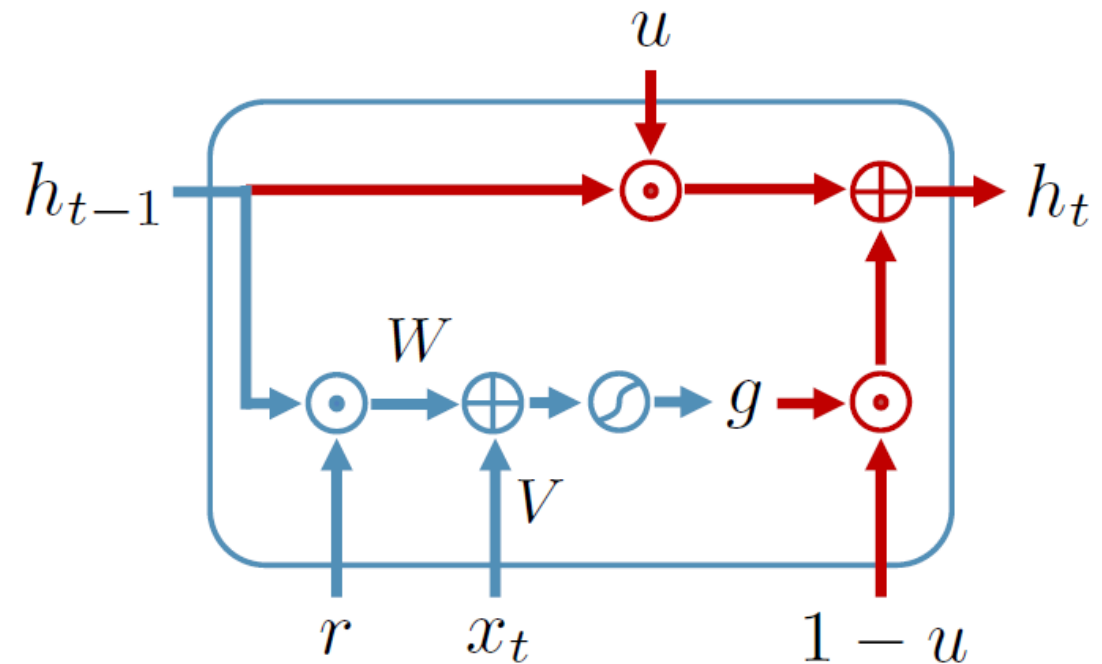
$$u = \sigma(W h_{t-1} + V x_t + b_u)$$

$$r = \sigma(W h_{t-1} + V x_t + b_r)$$

$$g = \text{Tanh}(W(h_{t-1} \odot r) + V x_t + b_g)$$

$$h_t = (1 - u) \odot g + u \odot h_{t-1}$$

GRU





LSTM or GRU

❖ LSTM

- more parameters => longer training
- More flexible

❖ GRU

- Less parameters => faster training

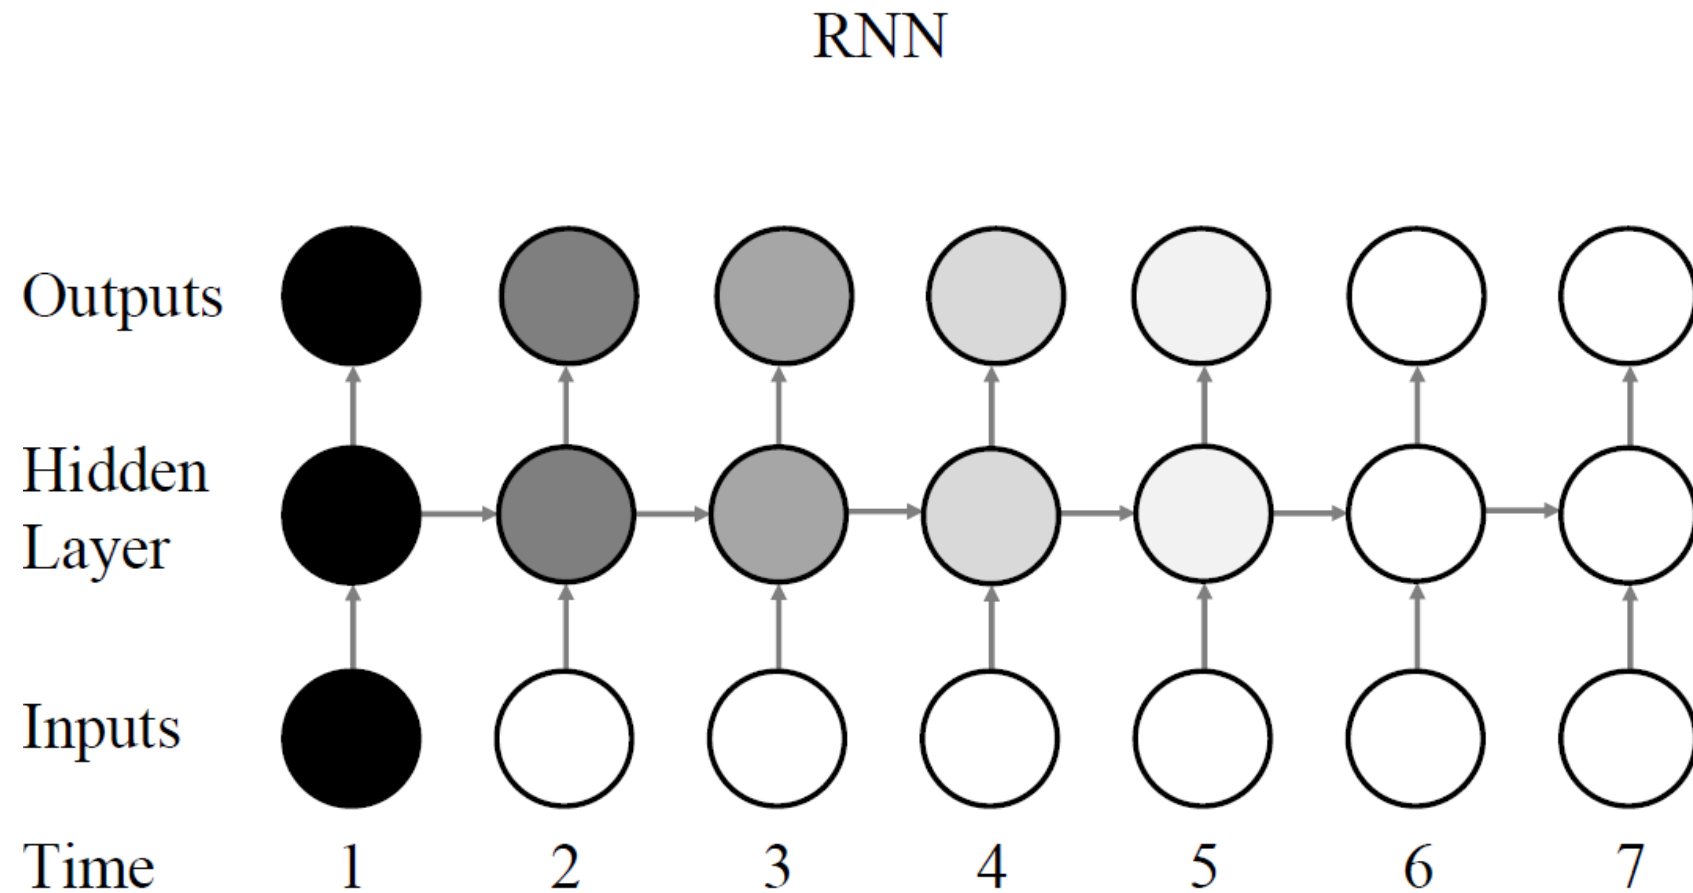
❖ Train using LSTM first

❖ Train using GRU next

❖ Choose GRU if performance is similar

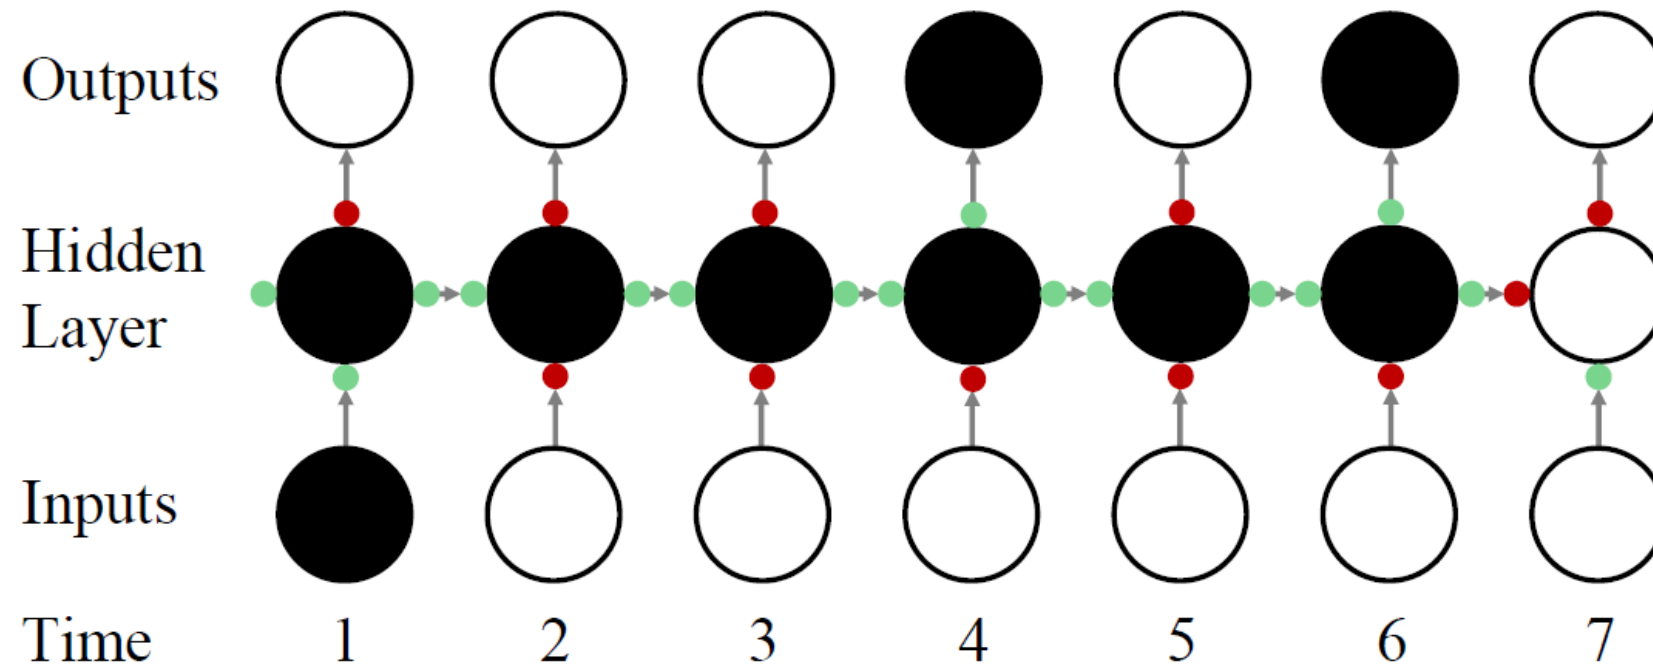


RNN Information Flow





LSTM Information Flow





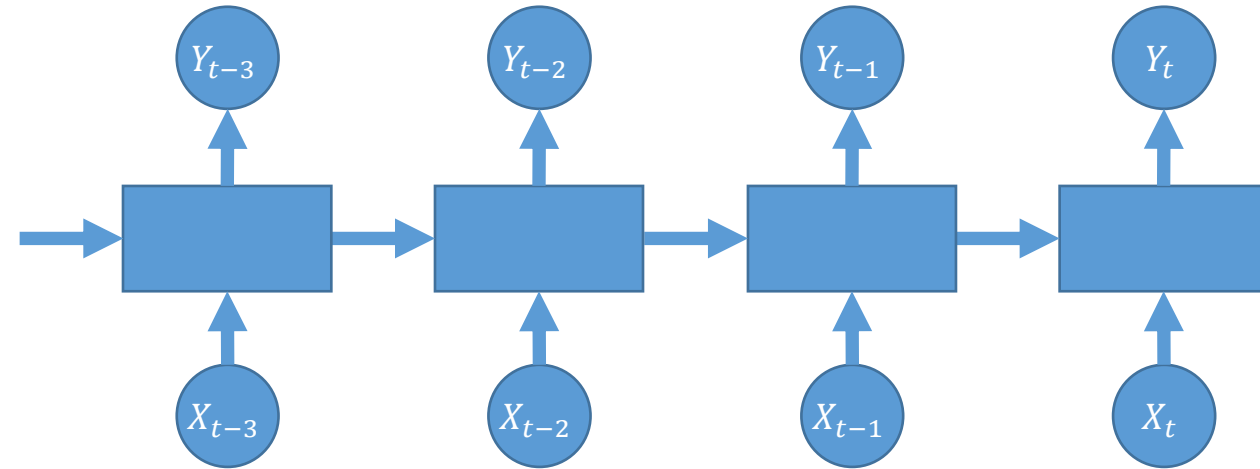
Regularization

- ❖ L2 regularization is very effective
- ❖ Dropout can be applied to V , U but not W (memory)



Text and Language Modeling

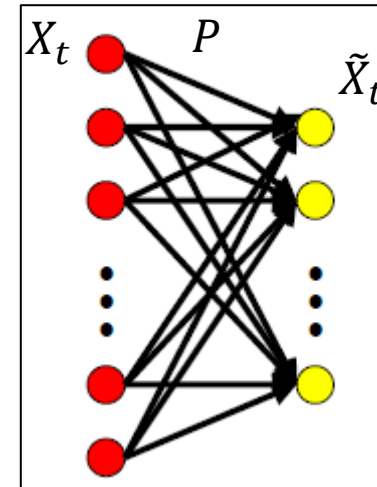
- ❖ Input as one-hot vector
- ❖ For text:
 - each letter is a vector
 - number of characters = dimension of vector
 - Includes upper case, lower case, hyphenated, commas, apostrophes as characters
 - Output: predict next character/word
- ❖ For language:
 - each word is a vector
 - Dictionary of all inputs
 - size of dictionary = dimension of vector
 - Includes upper case, lower case, hyphenated, commas, apostrophes as words
 - Output: predict next word/sentence





Curse of Dimensionality

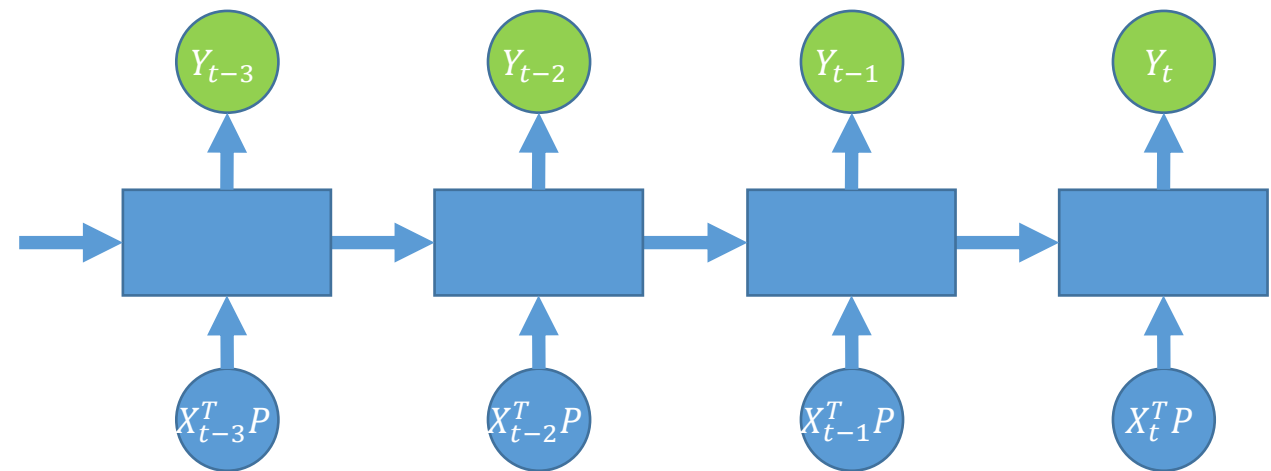
- ❖ Number of training samples = multiple of each dimension
- ❖ Observation: sparse space
 - Vertices of the space used not volume
 - Density = $N/2^N$
 - Highly inefficient
- ❖ Observation: vectors are unordered
 - Same length
- ❖ Idea: project to lower dimension space
 - Input vector is X_t : $1 \times N$
 - Projection is $1 \times M$, $M < N$
 - Projection function P : $M \times N$
 - Projection: $\tilde{X}_t = PX_t^T$
 - Learn the projection function, P
 - Unsupervised





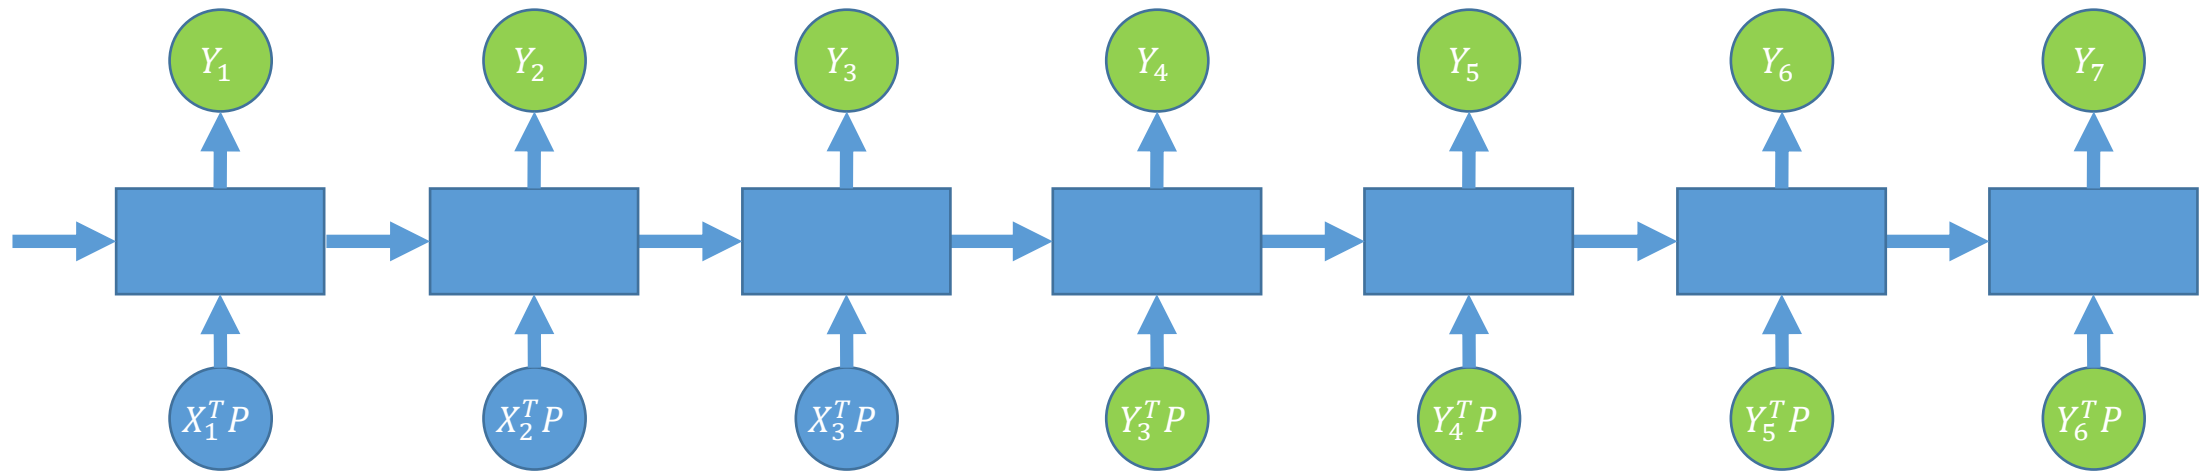
Language Synthesis: Training

❖ Use BPTT to train model





Language Synthesis: Generation



- ❖ Provide first few inputs
- ❖ Let the network feedback output back in



Beam Search

❖ Output: softmax gives most likely next character (or word)

- Probability distribution over all dictionary vectors
- Greedy: susceptible to propagating errors

❖ Beam search:

- pick a number of non-max outcomes (aka hypothesis)
- Evaluate each hypothesis by its overall probability: $\prod_i p(y_i)$
- Prune weak hypothesis at each iteration
- Repeat

The quick br

$$\begin{bmatrix} a = 0.3 \\ b = 0.01 \\ c = 0.01 \\ \vdots \\ o = 0.4 \\ \vdots \end{bmatrix}$$



The quick br

$$\begin{bmatrix} a = 0.3 \\ b = 0.01 \\ c = 0.01 \\ \vdots \\ o = 0.4 \\ \vdots \end{bmatrix}$$

Arrows point from the highlighted $a = 0.3$ and $o = 0.4$ to the following vectors:

$$\begin{bmatrix} a = 0.01 \\ b = 0.1 \\ c = 0.3 \\ \vdots \end{bmatrix}$$

$$\begin{bmatrix} a = 0.2 \\ b = 0.01 \\ \vdots \\ w = 0.6 \\ \vdots \end{bmatrix}$$

$$\begin{bmatrix} \text{braa} = 0.3 \times 0.01 \\ \text{brab} = 0.3 \times 0.1 \\ \text{brac} = 0.3 \times 0.3 \\ \vdots \\ \text{broa} = 0.4 \times 0.2 \\ \text{brob} = 0.4 \times 0.01 \\ \text{brow} = 0.4 \times 0.6 \\ \vdots \end{bmatrix}$$



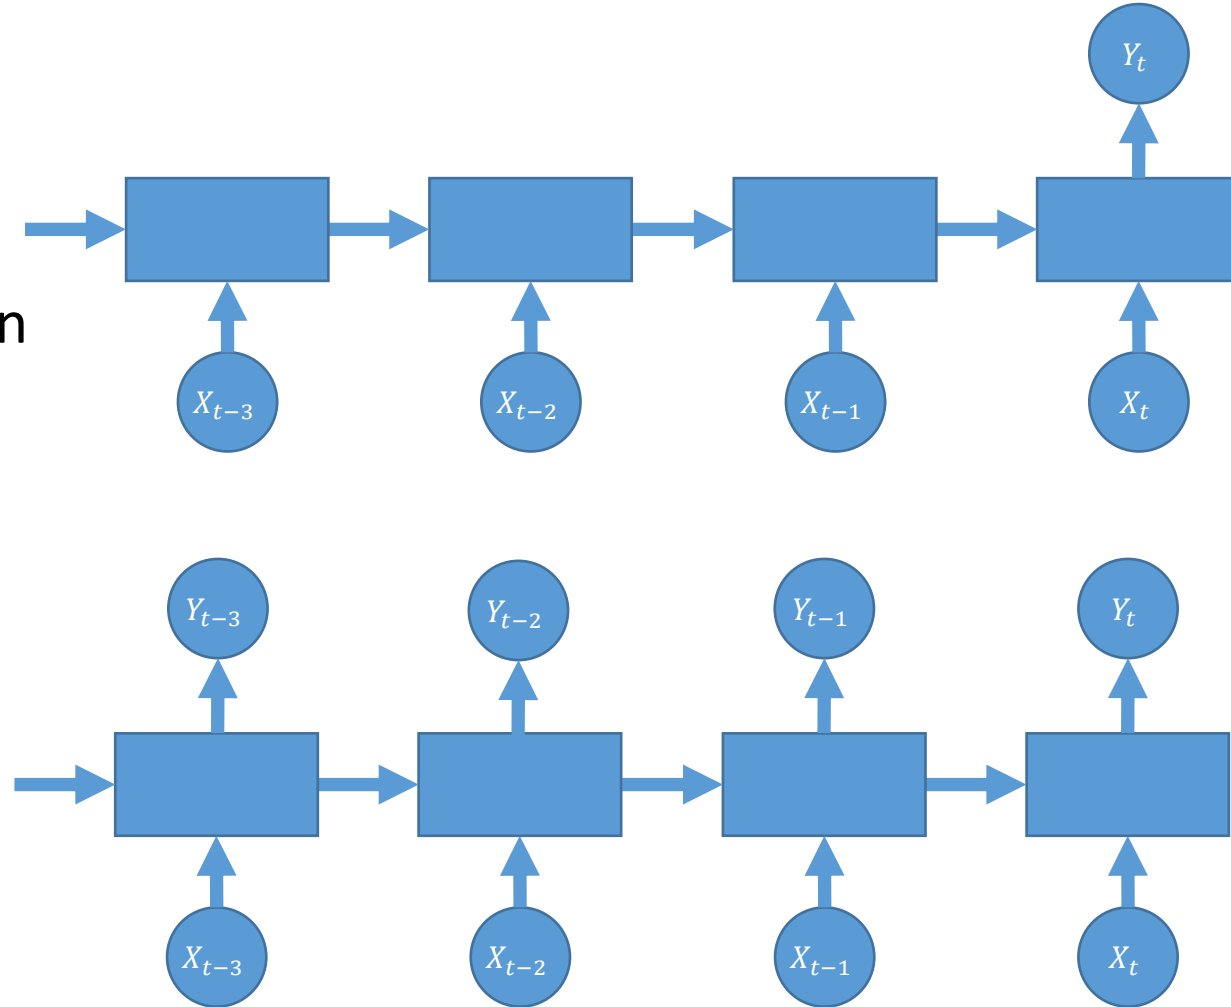
Sampled Softmax

- ❖ Regular softmax: $e^{y_j} / \sum_i e^{y_i}$
 - j is the target output
 - i is the negative (non-targets)
- ❖ Problem: possible outputs i is large
 - large dictionary
- ❖ Randomly select negatives and use in softmax
 - Importance sampling (speedup x19)
 - Adaptive importance sampling (x100)
 - Target sampling (AIS but partitioning training data to limit words)
- ❖ Alternatives:
 - Self normalization (x15 higher accuracy)
 - Noise contrastive estimator (x45 higher accuracy)



Phenome Recognition

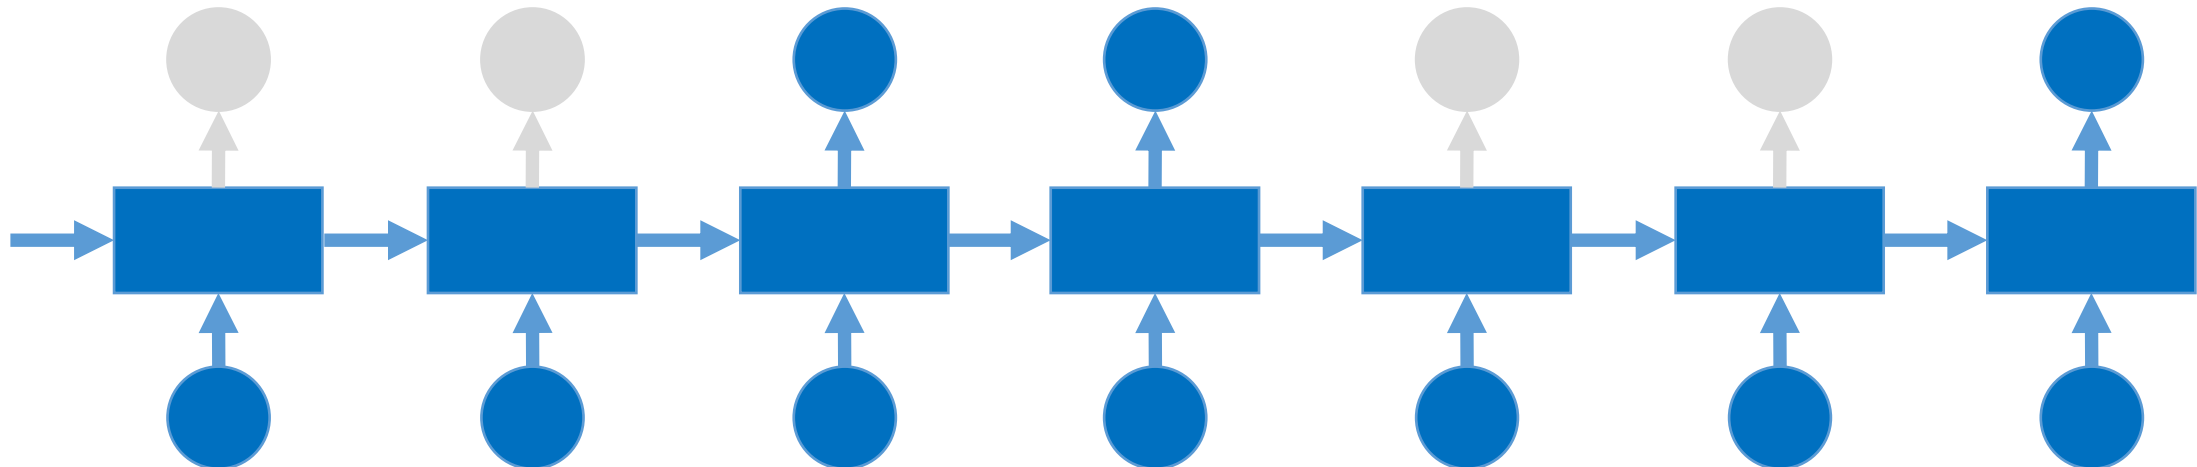
- ❖ Input: sequence of spectral data
- ❖ Output: phenome
- ❖ In reality output is produced in each iteration
 - Ignored until end
- ❖ Training: consider error at each iteration





Speech Recognition

- ❖ Input: sequence of inputs data (spectra)
- ❖ Output: **asynchronously** sequence of symbols (phenomes)





Speech Recognition Training

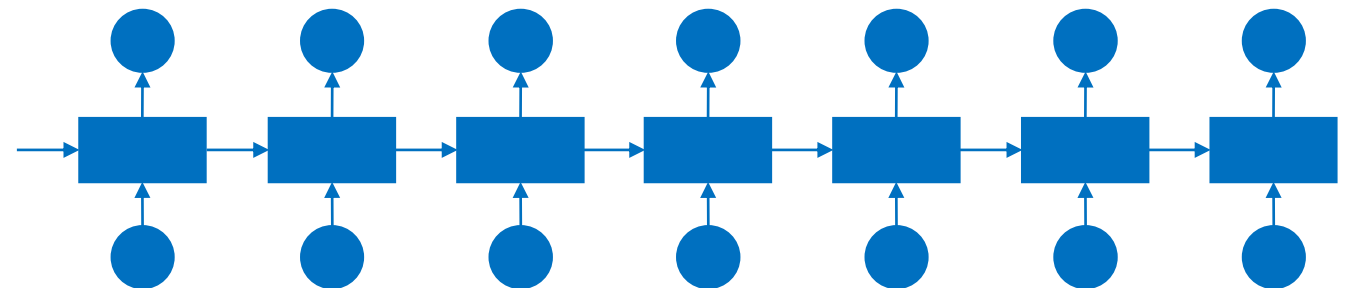
- ❖ No timing information
- ❖ Problem 1: output is a probability distribution over all symbols (phenomes)
- ❖ Problem 2: Can't differentiate between symbol repetition and symbol extension
- ❖ Problem 3: Even if the sequence is known, the timed output is not



Speech Recognition Training

- ❖ No timing information
- ❖ Problem 1: output is a probability distribution over all symbols (phenomes)
- ❖ Solution: merge the symbols

a	0.1	0.15	0.4	0.55	0.45	0.2	0.15
b	0.5	0.6	0.3	0.05	0.1	0.05	0.15
d	0.1	0.05	0.2	0.05	0.3	0.3	0.25
e	0.1	0.10	0.05	0.1	0.05	0.2	0.15
i	0.1	0.05	0.05	0.2	0.05	0.1	0.2
f	0.1	0.05	0.05	.05	0.05	0.05	0.2





Speech Recognition Training

❖ Problem 2: Can't differentiate between symbol repetition and symbol extension

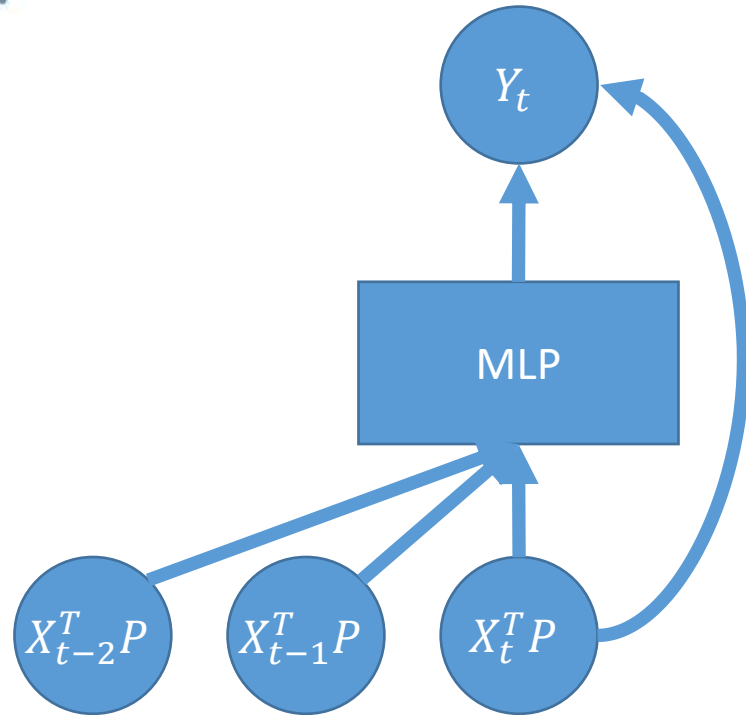


Speech Recognition Training

- ❖ No timing information
- ❖ Problem 1: output is a probability distribution over all symbols (phenomes)
- ❖ Problem 2: Can't differentiate between symbol repetition and symbol extension
- ❖ Problem 3: Even if the sequence is known, the timed output is not



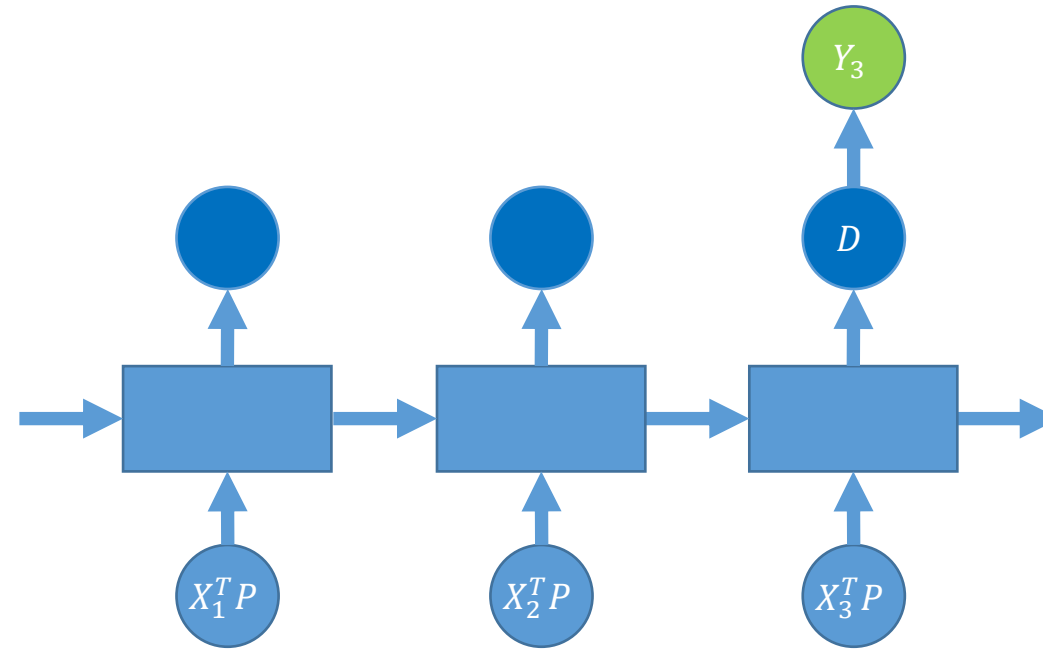
TDNN Model



❖ Predict characters/words based on last N



Language Synthesis: Generation



- ❖ Provide first few inputs
- ❖ After last input, generate a probability distribution over all dictionary entries
- ❖ Draw an entry from the dictionary with the highest probability