

Convolution Neural Nets

Part 1: Forward Propagation

Convolution:

Simple Convolution: $H \times W$ image with $h \times h$ filter: One matrix convolved by another matrix. Output is $H \times W$

```
def convHWpad(im,f):  
  
    H,W = im.shape  
    h,h = f.shape          #filter  
    hh = h//2  
    im_pad = np.zeros((H+hh*2,W+hh*2))  
    im_pad[hh:H+hh,hh:W+hh] = im  
    out = np.zeros((H,W))  
  
    for r in range(W):  
        for c in range(H):  
            prod = im_pad[r:r+h, c:c+h]*f  
            out[r,c] = np.sum(prod,axis=(0,1))  
    return out
```

Convolution: $H \times W \times C$ image with $h \times h \times C$ filter: One image with *multiple* channels convolved with a filter of equal channels. Output is $H \times W$

```
def convHWCpad(im,f):  
  
    H,W,C = im.shape  
    h,h,C = f.shape          #filter  
    hh = h//2  
    im_pad = np.zeros((H+hh*2,W+hh*2,C))  
    im_pad[hh:H+hh,hh:W+hh,:] = im  
    out = np.zeros((H,W))  
  
    for r in range(W):  
        for c in range(H):  
            prod = im_pad[r:r+h, c:c+h,:]*f  
            out[r,c] = np.sum(prod,axis=(0,1,2))  
    return out
```

Convolution: $H \times W \times C$ image with $n \times h \times h \times C$ filter: One image with multiple channels convolved with *multiple* filters of equal channels. Output $H \times W \times n$

```
def convHWNpad(im,f):
```

```
    H,W,C = im.shape
    n,h,h,C = f.shape          #filter
    hh = h//2
    im_pad = np.zeros((H+hh*2,W+hh*2,C))
    im_pad[hh:H+hh,hh:W+hh,:] = im
    out = np.zeros((H,W,n))

    for j in range(n):
        for r in range(W):
            for c in range(H):
                prod = im_pad[r:r+h, c:c+h,:]*f[j]
                out[r,c,j] = np.sum(prod,axis=(0,1,2))
    return out
```

Convolution: $N \times H \times W \times C$ image with $n \times h \times h \times C$ filter: *Multiple* images (batch/sample size) with multiple channels convolved with *multiple* filters of equal channels. Output is $N \times H \times W \times n$

```
def convNHWNpad(im,f):
```

```
    N,H,W,C = im.shape
    n,h,h,C = f.shape          #filter
    hh = h//2
    im_pad = np.zeros((N,H+hh*2,W+hh*2,C))
    im_pad[:,hh:H+hh,hh:W+hh,:] = im
    out = np.zeros((N,H,W,n))

    for i in range(N):
        for j in range(n):
            for r in range(W):
                for c in range(H):
                    prod = im_pad[i,r:r+h, c:c+h,:]*f[j]
                    out[i,r,c,j] = np.sum(prod,axis=(0,1,2))
    return out
```

Max-Pooling:

Simple max-pooling: $H \times W$ image with 2×2 filter, stride=2: One matrix, max is determined in 2×2 , stride = 2. Output is $W/2 \times H/2$.

```
def maxpoolHW(im):
    H,W = im.shape
    out = np.zeros((H//2,W//2))
    for r in range(0,H,2):
        for c in range(0,W,2):
            out[r//2,c//2] = np.max(im[r:r+2,c:c+2])
    return out
```

max-pooling: $H \times W \times C$ image with 2×2 filter, stride=2: multiple channels, max is determined in 2×2 per channel, stride = 2. Output is $W/2 \times H/2 \times C$.

```
def maxpoolHWC(im):
    H,W,C = im.shape
    out = np.zeros((H//2,W//2,C))
    for r in range(0,H,2):
        for c in range(0,W,2):
            for j in range(C):
                out[r//2,c//2,j] = np.max(im[r:r+2,c:c+2,j])
    return out
```

max-pooling: $N \times H \times W \times C$ image with 2×2 filter, stride=2: multiple channels, max is determined in 2×2 per channel, stride = 2. Output is $N \times W/2 \times H/2 \times C$.

```
def maxpoolNHWC(im):
    N,H,W,C = im.shape
    out = np.zeros((N,H//2,W//2,C))
    for i in range(N):
        for r in range(0,H,2):
            for c in range(0,W,2):
                for j in range(C):
                    out[i,r//2,c//2,j] = np.max(im[i,r:r+2,c:c+2,j])
    return out
```

ReLU Activation:

Simple ReLU on a $H \times W$ image: blocks negative values of the image. Output is $H \times W$.

```
def ReLU(im):
    return (img>0)*img
```

ReLU on a $N \times H \times W \times C$ image: blocks negative values of the image. Output is $N \times H \times W \times C$.

```
def ReLU(im):
    return (img>0)*img
```

Flatten:

Simple flattening of $H \times W$ image: Output is $1 \times HW$.

```
def flatten(im):  
    H,W = im.shape  
    return img.reshape(H*W)
```

Flattening of $H \times W \times C$ image: Output is $1 \times HWC$.

```
def flatten(im):  
    H,W,C = im.shape  
    return img.reshape(H*W*C)
```

Simple flattening of $N \times H \times W \times C$ image: Output is $N \times HWC$.

```
def flatten(im):  
    N,H,W,C = im.shape  
    return img.reshape(N,H*W*C)
```

Part 2: Backward Propagation

Backprop Theory

Forward Layers: Input \rightarrow $\underbrace{\text{Convolution} \rightarrow \text{ReLU} \rightarrow \text{Max-pooling}}_{\text{repeats}}$ \rightarrow Flatten \rightarrow $\underbrace{\text{Sigmoid}}_{\text{repeats}}$ \rightarrow Softmax

\mathcal{L} :	loss function
l :	layer number (L = last layer)
W^l :	Weight from layer l
a^l :	activation function in layer l
Z^l :	Input to activation functions
$\frac{\partial a^l}{\partial z^l} = g^l$:	derivative of the activation function

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial W^{L-2}} &= \frac{\partial \mathcal{L}}{\partial a^L} \left(\frac{\partial a^L}{\partial Z^L} \right) \frac{\partial Z^L}{\partial a^{L-1}} \left(\frac{\partial a^{L-1}}{\partial Z^{L-1}} \right) \frac{\partial Z^{L-1}}{\partial a^{L-2}} \left(\frac{\partial a^{L-2}}{\partial Z^{L-2}} \right) \frac{\partial Z^{L-2}}{\partial W^{L-2}} \\ &= \frac{\partial \mathcal{L}}{\partial a^L} \quad (g^L) \quad W^{L-1} \quad (g^{L-1}) \quad W^{L-2} \quad (g^{L-2}) \quad a^{L-2} \\ \Rightarrow \frac{\partial \mathcal{L}}{\partial a^l} &= \frac{\partial \mathcal{L}}{\partial a^{l+1}} (g^l) W^l, \quad \frac{\partial \mathcal{L}}{\partial W^l} = \frac{\partial \mathcal{L}}{\partial a^l} a^l, \quad \frac{\partial \mathcal{L}}{\partial b^l} = \frac{\partial \mathcal{L}}{\partial a^l} \end{aligned}$$

Summary:

For each layer calculate the gradient of activation function g^{l+1} , then the partial derivative of the loss wrt the activation function, $\frac{\partial \mathcal{L}}{\partial a^l}$, then partial derivative of loss wrt the weights, $\frac{\partial \mathcal{L}}{\partial w^l}$.

Backprop Layer:

```
def backproplayer(a_l, da_lp1, g_l, W_l):
    da_l = np.dot(da_lp1, W_l.T) * g_l
    dW_l = np.dot(a_l.T, da_l)
    db_l = np.sum(da_l, axis=0, keepdims=True)
    return da_l, dW_l, db_l
```

Backward Softmax (L):

$$g_{\text{softmax}} = \frac{\partial a_{\text{softmax}}}{\partial Z} = 1$$

$$\frac{\partial \mathcal{L}}{\partial a^L} = \frac{\partial \mathcal{L}}{\partial a_{\text{softmax}}} = \delta \left(\frac{\partial a_{\text{softmax}}}{\partial Z} \right) = \delta$$

$$\frac{\partial \mathcal{L}}{\partial W^{L-1}} = \delta (a^{L-1})$$

Backprop Layer:

```
da_Lm1, dW_L, db_L = backproplayer(softmax(a_Lm1,W_L,b_L), delta, 1, W_L):
```

Backward Sigmoid ($L - 1$):

$$g^{L-1} = \frac{\partial a_{\text{sigmoid}}^{L-1}}{\partial Z^{L-1}} = a_{\text{sigmoid}}^{L-1} (1 - a_{\text{sigmoid}}^{L-1})$$

$$\frac{\partial \mathcal{L}}{\partial a^{L-1}} = \frac{\partial \mathcal{L}}{\partial a^L} (g^{L-1}) W^L = \delta (g^{L-1})$$

$$\frac{\partial \mathcal{L}}{\partial W^{L-1}} = \frac{\partial \mathcal{L}}{\partial a^{L-1}} a_{\text{sigmoid}}^{L-2}$$

Backward Sigmoid (l):

$$g^l = \frac{\partial a_{\text{sigmoid}}^l}{\partial Z^l} = a_{\text{sigmoid}}^l (1 - a_{\text{sigmoid}}^l)$$

$$\frac{\partial \mathcal{L}}{\partial a^l} = \frac{\partial \mathcal{L}}{\partial a^{l+1}} (g^l) W^{l+1}$$

$$\frac{\partial \mathcal{L}}{\partial W^l} = \frac{\partial \mathcal{L}}{\partial a^l} a^{l-1}$$

Backward Flatten (l):

Flattening is a re-organization of the feature maps from the previous layer (max-pooling layer). It doesn't change the values in the previous layer, i.e. $\frac{\partial a_{\text{flatten}}}{\partial Z} = 1$. There are no weights connecting the previous layer to the flattened layer.

$$g^l = 1$$

$$\frac{\partial \mathcal{L}}{\partial a^l} = \frac{\partial \mathcal{L}}{\partial a^{l+1}} W^l$$

Simple unflattening of $N \times HWC$ image: Output is $N \times H \times W \times C$.

```
def unflatten(a,N,H,W,C):  
    return a.reshape(N,H,W,C)
```

Backward Max-pool (l):

Max-pooling selects the max in a 2x2. So for the values that match the max, $g = 1$, otherwise it is 0. There are no weights connecting the previous layer (convolution layer) to the max-pooling.

1	2	3	4
8	7	5	6
9	11	12	10
10	7	13	9

z^l

8	6
11	13

a^l

0	0	0	0
1	0	0	1
0	1	0	0
0	0	1	0

g^l

$$g^l = \begin{cases} 1 & z^l = a^l \\ 0 & \text{otherwise} \end{cases}$$

$$\frac{\partial \mathcal{L}}{\partial a^l} = \frac{\partial \mathcal{L}}{\partial a^{l+1}} W^{l+1}$$

Simple Gradient of Max-pooling: Assumes input is HxW and max pooling was 2x2, stride = 2

```
def g_maxpool(z_l, a_mp):
    g = np.repeat(a_mp, 2, axis=1)
    g = np.repeat(g, 2, axis=0)
    g = (g == z_l) * 1.
    return g
```

#a_mp is the output of maxpooling
#duplicate the columns
#duplicate the rows
#z_l is the input to max-pooling

Backward Convolution:

$$\frac{dL}{da^l} = \text{convolution}\left(W, \frac{dL}{da^{l+1}}\right)$$

$$\frac{dL}{dW^l} = \text{convolution}\left(a^{l-1}, \frac{dL}{da^l}\right)$$

See <https://www.jefkine.com/general/2016/09/05/backpropagation-in-convolutional-neural-networks/>

Note that convolution used in convolutional NNet is not a mathematical convolution. In mathematical convolution you flip one of the inputs about its axis (in 2D you flip horizontally, then vertically, hence it's a rotation by 180 degrees).